

Used_Car_Listings_Features_and_Price_Prediction

June 15, 2023

1 Used Car Listings: Features and Price Prediction

- By John David

2 About Dataset

2.1 Dataset Description

This dataset contains information about used car listings obtained through data scraping from cars.com. The dataset includes various features of the cars such as brand, model, year, mileage, engine details, transmission type, fuel type, drivetrain, and more. Additionally, it provides details about specific features like adaptive cruise control, navigation system, power liftgate, backup camera, and others.

The purpose of this dataset is to facilitate the development and evaluation of machine learning models for car price prediction. The dataset can be used for tasks such as regression analysis, feature engineering, and model training.

Please note that the dataset is collected from a specific source and represents a snapshot of used car listings at a particular point in time. The dataset may contain missing values and some discrepancies. It is recommended to perform data preprocessing and exploratory analysis before using it for research or predictive modeling.

We hope that this dataset serves as a valuable resource for the data science community and contributes to the advancement of automotive analytics and pricing models.

3 Feature Descriptions

- brand: Brand of the car.
- model: Model of the car.
- year: Year of production of the car.
- mileage: Mileage of the car.
- engine: Information about the car's engine.
- engine_size: Size of the car's engine.
- transmission: Type of transmission of the car.
- automatic_transmission: A binary value indicating the presence of automatic transmission (1: Yes, 0: No).

- fuel_type: Fuel type of the car.
- drivetrain: Drivetrain type of the car.
- min_mpg: Minimum fuel efficiency of the car.
- max_mpg: Maximum fuel efficiency of the car.
- damaged: A binary value indicating the presence of damage in the car (1: Yes, 0: No).
- first_owner: Is the car a 1-owner vehicle? (1: Yes, 0: No).
- personal_using: Is the car for personal use only? (1: Yes, 0: No).
- turbo: A binary value indicating the presence of a turbocharger in the car (1: Yes, 0: No).
- alloy_wheels: Are there alloy wheels on the car? (1: Yes, 0: No).
- adaptive_cruise_control: A binary value indicating the presence of adaptive cruise control (1: Yes, 0: No).
- navigation_system: A binary value indicating the presence of a navigation system (1: Yes, 0: No).
- power_liftgate: A binary value indicating the presence of a power liftgate (1: Yes, 0: No).
- backup_camera: A binary value indicating the presence of a backup camera (1: Yes, 0: No).
- keyless_start: A binary value indicating the presence of keyless start system (1: Yes, 0: No).
- remote_start: A binary value indicating the presence of a remote start system (1: Yes, 0: No).
- sunroof/moonroof: A binary value indicating the presence of a sunroof/moonroof (1: Yes, 0: No).
- automatic_emergency_braking: A binary value indicating the presence of automatic emergency braking system (1: Yes, 0: No).
- stability_control: A binary value indicating the presence of stability control system (1: Yes, 0: No).
- leather_seats: Are there leather seats in the car? (1: Yes, 0: No).
- memory_seat: Are there memory seats in the car? (1: Yes, 0: No).
- third_row_seating: A binary value indicating the presence of third row seating (1: Yes, 0: No).
- apple_car_play/android_auto: A binary value indicating the presence of Apple CarPlay / Android Auto integration (1: Yes, 0: No).
- bluetooth: A binary value indicating the presence of Bluetooth connectivity (1: Yes, 0: No).
- usb_port: A binary value indicating the presence of USB ports (1: Yes, 0: No).
- heated_seats: Are there heated seats in the car? (1: Yes, 0: No).
- interior_color: Interior color of the car.
- exterior_color: Exterior color of the car.
- price: Price of the car. This feature is the target feature of this dataset.

4 Importing the libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set_style('darkgrid')
plt.style.use('bmh')
import warnings
```

```
warnings.filterwarnings('ignore')
import os
```

5 Loading the dataset

```
[2]: # Train dataset
train = pd.read_csv('/content/train.csv')
test = pd.read_csv('/content/test.csv')
print('TRAIN DATASET SHAPE: ', train.shape)
print('TEST DATASET SHAPE: ', test.shape)
```

```
TRAIN DATASET SHAPE: (19109, 36)
TEST DATASET SHAPE: (4778, 36)
```

5.1 Train dataset

```
[3]: # train dataset view
train.head()
```

```
[3]:
```

	brand	model	year	mileage	\
0	Mazda	CX-9 Touring	2022.0	6580.0	
1	Alfa Romeo	Stelvio 4DR SUV AWD	2018.0	62717.0	
2	Chevrolet	Silverado 1500 Limited LT	2022.0	15463.0	
3	Land Rover	Range Rover P400 SE	2022.0	1281.0	
4	Mitsubishi	Outlander Sport 2.0 SE	2021.0	36727.0	

	engine	engine_size	transmission	\
0	2.5L I4 16V GDI DOHC Turbo	2.5	6-Speed Automatic	
1	Twin Turbo Premium Unleaded V-6 2.9 L/176	2.9	Automatic	
2	5.3L V8 16V GDI OHV	5.3	Automatic	
3	3.0L I6 24V GDI DOHC Turbo	3.0	Automatic	
4	2.0L I4 16V MPFI DOHC	2.0	Automatic CVT	

	automatic_transmission	fuel_type	drivetrain	...	leather_seats	\
0	1.0	Gasoline	Four-wheel Drive	...	1.0	
1	1.0	Gasoline	Four-wheel Drive	...	0.0	
2	1.0	Gasoline	Four-wheel Drive	...	0.0	
3	1.0	Gasoline	Four-wheel Drive	...	1.0	
4	1.0	Gasoline	Front-wheel Drive	...	0.0	

	memory_seat	third_row_seating	apple_car_play/android_auto	bluetooth	\
0	0.0	1.0	1.0	0.0	
1	1.0	0.0	1.0	1.0	
2	0.0	0.0	0.0	0.0	
3	1.0	0.0	0.0	0.0	
4	0.0	0.0	1.0	1.0	

	usb_port	heated_seats	interior_color		exterior_color	price
0	0.0	1.0	Black		Snowflake White Pearl	36789
1	0.0	0.0	Black	Rosso	Competizione Tri-Coat	39993
2	0.0	1.0	Jet Black		Silver Ice Metallic	46986
3	0.0	1.0	Ebony		Fuji White	141999
4	1.0	1.0	Black		White	21595

[5 rows x 36 columns]

6 Data preprocessing 1

```
[4]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19109 entries, 0 to 19108
Data columns (total 36 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   brand                                19107 non-null  object
1   model                                19107 non-null  object
2   year                                19107 non-null  float64
3   mileage                              19107 non-null  float64
4   engine                               19062 non-null  object
5   engine_size                          17860 non-null  float64
6   transmission                         19004 non-null  object
7   automatic_transmission               19107 non-null  float64
8   fuel_type                            19107 non-null  object
9   drivetrain                           19107 non-null  object
10  min_mpg                              16103 non-null  float64
11  max_mpg                              16103 non-null  float64
12  damaged                              18940 non-null  float64
13  first_owner                          18803 non-null  float64
14  personal_using                       18920 non-null  float64
15  turbo                                19107 non-null  float64
16  alloy_wheels                         19107 non-null  float64
17  adaptive_cruise_control              19107 non-null  float64
18  navigation_system                    19107 non-null  float64
19  power_liftgate                       19107 non-null  float64
20  backup_camera                        19107 non-null  float64
21  keyless_start                        19107 non-null  float64
22  remote_start                         19107 non-null  float64
23  sunroof/moonroof                     19107 non-null  float64
24  automatic_emergency_braking          19107 non-null  float64
25  stability_control                    19107 non-null  float64
26  leather_seats                        19107 non-null  float64
```

```

27 memory_seat          19107 non-null float64
28 third_row_seating    19107 non-null float64
29 apple_car_play/android_auto 19107 non-null float64
30 bluetooth            19107 non-null float64
31 usb_port             19107 non-null float64
32 heated_seats         19107 non-null float64
33 interior_color       17952 non-null object
34 exterior_color       18900 non-null object
35 price                19107 non-null object

```

dtypes: float64(27), object(9)

memory usage: 5.2+ MB

6.1 Test dataset

```
[5]: # test dataset
test.head()
```

```
[5]:
   brand      model  year  mileage \
0  FIAT    124 Spider Lusso  2020.0  30830.0
1  Lexus    RX 450h F Sport  2019.0  58274.0
2  Kia      Stinger GT1    2018.0  69303.0
3  MINI  Convertible Cooper S  2023.0   5026.0
4  Nissan      Murano S    2018.0  61180.0

   engine      engine_size  transmission \
0  1.4L I4 16V MPFI SOHC Turbo      1.4 6-Speed Automatic
1  3.5L V6 24V PDI DOHC Hybrid      3.5  Automatic CVT
2  3.3L V6 24V GDI DOHC Twin Turbo    3.3 8-Speed Automatic
3  2.0L I4 16V GDI DOHC Turbo      2.0  Automatic
4  3.5L V6 24V MPFI DOHC      3.5  Automatic CVT

   automatic_transmission fuel_type  drivetrain  ...  leather_seats \
0                1.0  Gasoline  Rear-wheel Drive  ...            1.0
1                1.0   Hybrid  Four-wheel Drive  ...            1.0
2                1.0  Gasoline  Four-wheel Drive  ...            1.0
3                1.0  Gasoline  Front-wheel Drive  ...             0.0
4                1.0  Gasoline  Four-wheel Drive  ...             0.0

   memory_seat  third_row_seating  apple_car_play/android_auto  bluetooth \
0            0.0                0.0                        0.0            1.0
1            1.0                0.0                        0.0            1.0
2            1.0                0.0                        1.0            1.0
3            0.0                0.0                        1.0            0.0
4            0.0                0.0                        1.0            1.0

   usb_port  heated_seats  interior_color      exterior_color  price
0         1.0           1.0         Saddle  Forte Black Metallic  26500

```

1	0.0	1.0	Rioja Red	Nightfall Mica	36987
2	0.0	1.0	Black	HiChroma Red	26997
3	0.0	1.0	Carbon Black	Chili Pepper Red	37923
4	1.0	0.0	Graphite	Gun Metallic	20490

[5 rows x 36 columns]

```
[6]: test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4778 entries, 0 to 4777
Data columns (total 36 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   brand                                4777 non-null   object
1   model                                4777 non-null   object
2   year                                4777 non-null   float64
3   mileage                              4777 non-null   float64
4   engine                               4771 non-null   object
5   engine_size                          4493 non-null   float64
6   transmission                         4760 non-null   object
7   automatic_transmission               4777 non-null   float64
8   fuel_type                            4777 non-null   object
9   drivetrain                           4777 non-null   object
10  min_mpg                              4032 non-null   float64
11  max_mpg                              4032 non-null   float64
12  damaged                              4723 non-null   float64
13  first_owner                          4689 non-null   float64
14  personal_using                       4721 non-null   float64
15  turbo                                4777 non-null   float64
16  alloy_wheels                         4777 non-null   float64
17  adaptive_cruise_control              4777 non-null   float64
18  navigation_system                    4777 non-null   float64
19  power_liftgate                       4777 non-null   float64
20  backup_camera                        4777 non-null   float64
21  keyless_start                        4777 non-null   float64
22  remote_start                         4777 non-null   float64
23  sunroof/moonroof                     4777 non-null   float64
24  automatic_emergency_braking          4777 non-null   float64
25  stability_control                    4777 non-null   float64
26  leather_seats                        4777 non-null   float64
27  memory_seat                          4777 non-null   float64
28  third_row_seating                    4777 non-null   float64
29  apple_car_play/android_auto          4777 non-null   float64
30  bluetooth                            4777 non-null   float64
31  usb_port                             4777 non-null   float64
32  heated_seats                         4777 non-null   float64
```

```

33 interior_color          4457 non-null  object
34 exterior_color          4717 non-null  object
35 price                   4777 non-null  object
dtypes: float64(27), object(9)
memory usage: 1.3+ MB

```

```
[7]: df = pd.concat([train, test])
df.reset_index(drop=True, inplace=True)
```

```
[8]: df.nunique()
```

```

[8]: brand                25
model                   3789
year                    63
mileage                21313
engine                 1289
engine_size             64
transmission           205
automatic_transmission    2
fuel_type              10
drivetrain              5
min_mpg                55
max_mpg                58
damaged                 2
first_owner             2
personal_using          2
turbo                   2
alloy_wheels            2
adaptive_cruise_control  2
navigation_system       2
power_liftgate          2
backup_camera           2
keyless_start           2
remote_start            2
sunroof/moonroof        2
automatic_emergency_braking  2
stability_control       2
leather_seats           2
memory_seat             2
third_row_seating       2
apple_car_play/android_auto  2
bluetooth               2
usb_port                2
heated_seats            2
interior_color          1053
exterior_color          2135
price                  10656

```

dtype: int64

```
[9]: df.describe()
```

```
[9]:
```

	year	mileage	engine_size	automatic_transmission	\
count	23884.000000	23884.000000	22353.000000	23884.000000	
mean	2017.765701	48137.358776	2.979636	0.917518	
std	4.975234	38084.892846	5.777195	0.275103	
min	1953.000000	0.000000	0.000000	0.000000	
25%	2016.000000	19514.500000	2.000000	1.000000	
50%	2019.000000	39428.000000	2.500000	1.000000	
75%	2021.000000	68611.750000	3.500000	1.000000	
max	2024.000000	383614.000000	454.000000	1.000000	

	min_mpg	max_mpg	damaged	first_owner	personal_using	\
count	20135.000000	20135.000000	23663.000000	23492.000000	23641.000000	
mean	20.833126	27.820685	0.218654	0.518304	0.717905	
std	5.991664	6.487040	0.413342	0.499675	0.450029	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	17.000000	24.000000	0.000000	0.000000	0.000000	
50%	20.000000	28.000000	0.000000	1.000000	1.000000	
75%	24.000000	31.000000	0.000000	1.000000	1.000000	
max	89.000000	100.000000	1.000000	1.000000	1.000000	

	turbo	...	sunroof/moonroof	automatic_emergency_braking	\
count	23884.000000	...	23884.000000	23884.000000	
mean	0.404706	...	0.411489	0.173254	
std	0.490845	...	0.492114	0.378475	
min	0.000000	...	0.000000	0.000000	
25%	0.000000	...	0.000000	0.000000	
50%	0.000000	...	0.000000	0.000000	
75%	1.000000	...	1.000000	0.000000	
max	1.000000	...	1.000000	1.000000	

	stability_control	leather_seats	memory_seat	third_row_seating	\
count	23884.000000	23884.000000	23884.000000	23884.000000	
mean	0.949715	0.696491	0.483085	0.148677	
std	0.218536	0.459783	0.499724	0.355777	
min	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	0.000000	0.000000	0.000000	
50%	1.000000	1.000000	0.000000	0.000000	
75%	1.000000	1.000000	1.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	

	apple_car_play/android_auto	bluetooth	usb_port	heated_seats
count	23884.000000	23884.000000	23884.000000	23884.000000
mean	0.387163	0.869997	0.354296	0.483462

std	0.487112	0.336314	0.478309	0.499737
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	1.000000	0.000000	0.000000
50%	0.000000	1.000000	0.000000	0.000000
75%	1.000000	1.000000	1.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000

[8 rows x 27 columns]

[10]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23887 entries, 0 to 23886
Data columns (total 36 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   brand                                23884 non-null  object
1   model                                23884 non-null  object
2   year                                23884 non-null  float64
3   mileage                              23884 non-null  float64
4   engine                               23833 non-null  object
5   engine_size                          22353 non-null  float64
6   transmission                         23764 non-null  object
7   automatic_transmission               23884 non-null  float64
8   fuel_type                            23884 non-null  object
9   drivetrain                           23884 non-null  object
10  min_mpg                              20135 non-null  float64
11  max_mpg                              20135 non-null  float64
12  damaged                              23663 non-null  float64
13  first_owner                          23492 non-null  float64
14  personal_using                       23641 non-null  float64
15  turbo                                23884 non-null  float64
16  alloy_wheels                         23884 non-null  float64
17  adaptive_cruise_control              23884 non-null  float64
18  navigation_system                    23884 non-null  float64
19  power_liftgate                       23884 non-null  float64
20  backup_camera                        23884 non-null  float64
21  keyless_start                        23884 non-null  float64
22  remote_start                         23884 non-null  float64
23  sunroof/moonroof                     23884 non-null  float64
24  automatic_emergency_braking          23884 non-null  float64
25  stability_control                    23884 non-null  float64
26  leather_seats                        23884 non-null  float64
27  memory_seat                          23884 non-null  float64
28  third_row_seating                    23884 non-null  float64
29  apple_car_play/android_auto          23884 non-null  float64
30  bluetooth                            23884 non-null  float64
```

```

31  usb_port                23884 non-null float64
32  heated_seats            23884 non-null float64
33  interior_color          22409 non-null object
34  exterior_color          23617 non-null object
35  price                    23884 non-null object
dtypes: float64(27), object(9)
memory usage: 6.6+ MB

```

```
[11]: df.isnull().sum()
```


```

[11]: brand                3
      model                3
      year                3
      mileage              3
      engine              54
      engine_size         1534
      transmission        123
      automatic_transmission  3
      fuel_type           3
      drivetrain          3
      min_mpg             3752
      max_mpg             3752
      damaged             224
      first_owner         395
      personal_using       246
      turbo               3
      alloy_wheels         3
      adaptive_cruise_control  3
      navigation_system     3
      power_liftgate        3
      backup_camera         3
      keyless_start         3
      remote_start          3
      sunroof/moonroof      3
      automatic_emergency_braking  3
      stability_control     3
      leather_seats         3
      memory_seat          3
      third_row_seating     3
      apple_car_play/android_auto  3
      bluetooth            3
      usb_port             3
      heated_seats         3
      interior_color       1478
      exterior_color       270
      price                3
      dtype: int64

```

```
[12]: sns.heatmap(df.isnull(), cmap='Reds')
```

```
[12]: <Axes: >
```



```
[13]: df1 = df.dropna()
df1.reset_index(drop=True, inplace=True)
```

```
[14]: df1
```

```
[14]:
```

	brand	model	year	mileage \
0	Honda	Pilot Special Edition	2021.0	51299.0

1	Mercedes-Benz	GL-Class GL 450 4MATIC	2014.0	79279.0
2	Mitsubishi	Outlander Sport SE	2021.0	31976.0
3	Porsche	968	1994.0	92025.0
4	MINI	Hardtop Cooper	2016.0	43418.0
...
17463	Land	Rover Range Rover Evoque SE	2016.0	106438.0
17464	Mitsubishi	Outlander SE	2019.0	56555.0
17465	Nissan	GT-R Premium	2015.0	12091.0
17466	Maserati	Levante	2021.0	31187.0
17467	Kia	Optima Hybrid LX	2013.0	86261.0

	engine	engine_size \
0	3.5L V6 24V GDI SOHC	3.5
1	4.7L V8 32V GDI DOHC Twin Turbo	4.7
2	2.0L I4 16V MPFI DOHC	2.0
3	3.0L I4 16V MPFI DOHC	3.0
4	1.5L I3 12V GDI DOHC Turbo	1.5
...
17463	2.0L I4 16V GDI DOHC Turbo	2.0
17464	2.4L I-4 variable valve control, engine with 1...	2.4
17465	3.8L V-6 DOHC, variable valve control, twin tu...	3.8
17466	3.0L V6 24V GDI DOHC Twin Turbo	3.0
17467	2.4L I4 16V MPFI DOHC Hybrid	2.4

	transmission	automatic_transmission	fuel_type	drivetrain \
0	9-Speed Automatic	1.0	Gasoline	Four-wheel Drive
1	7-Speed Automatic	1.0	Gasoline	Four-wheel Drive
2	Automatic CVT	1.0	Gasoline	Four-wheel Drive
3	6-Speed Automatic	1.0	Gasoline	Rear-wheel Drive
4	6-Speed Automatic	1.0	Gasoline	Front-wheel Drive
...
17463	9-Speed Automatic	1.0	Gasoline	Four-wheel Drive
17464	Automatic	1.0	Gasoline	Four-wheel Drive
17465	Automatic	1.0	Gasoline	Four-wheel Drive
17466	8-Speed Automatic	1.0	Gasoline	Four-wheel Drive
17467	6-Speed Automatic	1.0	Hybrid	Front-wheel Drive

	leather_seats	memory_seat	third_row_seating \
0	1.0	1.0	1.0
1	1.0	1.0	1.0
2	0.0	0.0	0.0
3	1.0	0.0	0.0
4	1.0	0.0	0.0
...
17463	1.0	1.0	0.0
17464	0.0	0.0	1.0
17465	1.0	0.0	0.0

```

17466 ...      1.0      1.0      0.0
17467 ...      0.0      0.0      0.0

```

```

      apple_car_play/android_auto  bluetooth  usb_port  heated_seats  \
0                                1.0        1.0        0.0          1.0
1                                0.0        1.0        0.0          0.0
2                                0.0        1.0        0.0          0.0
3                                0.0        0.0        0.0          0.0
4                                0.0        1.0        1.0          0.0
...
17463 ...      0.0      1.0      1.0      1.0
17464 ...      1.0      1.0      1.0      1.0
17465 ...      0.0      1.0      0.0      1.0
17466 ...      0.0      1.0      0.0      1.0
17467 ...      0.0      1.0      0.0      0.0

```

```

      interior_color      exterior_color  price
0          Black  Modern Steel Metallic  35999
1          Black  Obsidian Black Metallic  22349
2          Gray          Silver  19499
3          Saddle          White  12960
4    Carbon Black  Electric Blue Metallic  16500
...
17463          Ebony  Santorini Black  16495
17464          Black  Mercury Gray Metallic  23988
17465          Black          Pearl White  92995
17466          Black          White  49850
17467    Sand / Black          Black  15595

```

[17468 rows x 36 columns]

7 Data Preprocessing 2

```
[15]: df1.isnull().sum()
```

```


[15]: brand          0
      model          0
      year          0
      mileage        0
      engine         0
      engine_size    0
      transmission   0
      automatic_transmission  0
      fuel_type      0
      drivetrain     0
      min_mpg        0

```

max_mpg	0
damaged	0
first_owner	0
personal_using	0
turbo	0
alloy_wheels	0
adaptive_cruise_control	0
navigation_system	0
power_liftgate	0
backup_camera	0
keyless_start	0
remote_start	0
sunroof/moonroof	0
automatic_emergency_braking	0
stability_control	0
leather_seats	0
memory_seat	0
third_row_seating	0
apple_car_play/android_auto	0
bluetooth	0
usb_port	0
heated_seats	0
interior_color	0
exterior_color	0
price	0
dtype:	int64

```
[16]: sns.heatmap(df1.isnull(), cmap='Blues')
```

```
[16]: <Axes: >
```



```
[17]: df1.shape
```

```
[17]: (17468, 36)
```

```
[18]: df1.nunique()
```

```
[18]: brand          25
      model        3115
      year          44
      mileage     16117
      engine        976
```

engine_size	57
transmission	149
automatic_transmission	2
fuel_type	7
drivetrain	4
min_mpg	53
max_mpg	55
damaged	2
first_owner	2
personal_using	2
turbo	2
alloy_wheels	2
adaptive_cruise_control	2
navigation_system	2
power_liftgate	2
backup_camera	2
keyless_start	2
remote_start	2
sunroof/moonroof	2
automatic_emergency_braking	2
stability_control	2
leather_seats	2
memory_seat	2
third_row_seating	2
apple_car_play/android_auto	2
bluetooth	2
usb_port	2
heated_seats	2
interior_color	905
exterior_color	1844
price	8410
dtype:	int64

```
[19]: # Checking the datatype
df1.dtypes
```

```
[19]: brand          object
      model          object
      year          float64
      mileage        float64
      engine          object
      engine_size     float64
      transmission    object
      automatic_transmission float64
      fuel_type        object
      drivetrain       object
      min_mpg          float64
```


max_mpg	float64
damaged	float64
first_owner	float64
personal_using	float64
turbo	float64
alloy_wheels	float64
adaptive_cruise_control	float64
navigation_system	float64
power_liftgate	float64
backup_camera	float64
keyless_start	float64
remote_start	float64
sunroof/moonroof	float64
automatic_emergency_braking	float64
stability_control	float64
leather_seats	float64
memory_seat	float64
third_row_seating	float64
apple_car_play/android_auto	float64
bluetooth	float64
usb_port	float64
heated_seats	float64
interior_color	object
exterior_color	object
price	object
dtype:	object

```
[20]: df1.columns
```

```
[20]: Index(['brand', 'model', 'year', 'mileage', 'engine', 'engine_size',
          'transmission', 'automatic_transmission', 'fuel_type', 'drivetrain',
          'min_mpg', 'max_mpg', 'damaged', 'first_owner', 'personal_using',
          'turbo', 'alloy_wheels', 'adaptive_cruise_control', 'navigation_system',
          'power_liftgate', 'backup_camera', 'keyless_start', 'remote_start',
          'sunroof/moonroof', 'automatic_emergency_braking', 'stability_control',
          'leather_seats', 'memory_seat', 'third_row_seating',
          'apple_car_play/android_auto', 'bluetooth', 'usb_port', 'heated_seats',
          'interior_color', 'exterior_color', 'price'],
          dtype='object')
```

```
[21]: df1.iloc[42]
```

```
[21]: brand          FIAT
      model          500 Sport
      year          2012.0
      mileage        79463.0
      engine        1.4L I4 16V MPFI SOHC
```

```

engine_size          1.4
transmission         6-Speed Automatic
automatic_transmission 1.0
fuel_type            Gasoline
drivetrain           Front-wheel Drive
min_mpg              30.0
max_mpg              38.0
damaged              0.0
first_owner          0.0
personal_using       0.0
turbo                0.0
alloy_wheels         1.0
adaptive_cruise_control 0.0
navigation_system    0.0
power_liftgate       0.0
backup_camera        0.0
keyless_start        0.0
remote_start         0.0
sunroof/moonroof     0.0
automatic_emergency_braking 0.0
stability_control    1.0
leather_seats        0.0
memory_seat          0.0
third_row_seating    0.0
apple_car_play/android_auto 0.0
bluetooth            1.0
usb_port             0.0
heated_seats         0.0
interior_color        Black
exterior_color        Rosso
price                ot Priced
Name: 42, dtype: object

```

```
[22]: df1['price'] = df1['price'].str.replace(',', '').str.replace('ot Priced','')
df.dtypes
```

```

[22]: brand          object
      model          object
      year          float64
      mileage        float64
      engine         object
      engine_size     float64
      transmission    object
      automatic_transmission float64
      fuel_type       object
      drivetrain      object
      min_mpg         float64

```

max_mpg	float64
damaged	float64
first_owner	float64
personal_using	float64
turbo	float64
alloy_wheels	float64
adaptive_cruise_control	float64
navigation_system	float64
power_liftgate	float64
backup_camera	float64
keyless_start	float64
remote_start	float64
sunroof/moonroof	float64
automatic_emergency_braking	float64
stability_control	float64
leather_seats	float64
memory_seat	float64
third_row_seating	float64
apple_car_play/android_auto	float64
bluetooth	float64
usb_port	float64
heated_seats	float64
interior_color	object
exterior_color	object
price	object
dtype:	object

```
[23]: df1['price'] = pd.to_numeric(df1['price'])
      df1['price'].dtype
```

```
[23]: dtype('float64')
```

```
[24]: df1.iloc[42]
```

```
[24]: brand          FIAT
      model          500 Sport
      year          2012.0
      mileage        79463.0
      engine         1.4L I4 16V MPFI SOHC
      engine_size          1.4
      transmission    6-Speed Automatic
      automatic_transmission  1.0
      fuel_type        Gasoline
      drivetrain       Front-wheel Drive
      min_mpg          30.0
      max_mpg          38.0
      damaged          0.0
```

first_owner	0.0
personal_using	0.0
turbo	0.0
alloy_wheels	1.0
adaptive_cruise_control	0.0
navigation_system	0.0
power_liftgate	0.0
backup_camera	0.0
keyless_start	0.0
remote_start	0.0
sunroof/moonroof	0.0
automatic_emergency_braking	0.0
stability_control	1.0
leather_seats	0.0
memory_seat	0.0
third_row_seating	0.0
apple_car_play/android_auto	0.0
bluetooth	1.0
usb_port	0.0
heated_seats	0.0
interior_color	Black
exterior_color	Rosso
price	NaN

Name: 42, dtype: object

```
[25]: df1.dropna(axis = 0, inplace=True)
```

```
[26]: df1.isnull().sum()
```

```
[26]: brand          0
      model          0
      year           0
      mileage        0
      engine         0
      engine_size    0
      transmission   0
      automatic_transmission  0
      fuel_type      0
      drivetrain     0
      min_mpg        0
      max_mpg        0
      damaged        0
      first_owner    0
      personal_using  0
      turbo          0
      alloy_wheels   0
      adaptive_cruise_control  0
```

```

navigation_system      0
power_liftgate         0
backup_camera          0
keyless_start          0
remote_start           0
sunroof/moonroof       0
automatic_emergency_braking 0
stability_control      0
leather_seats          0
memory_seat            0
third_row_seating      0
apple_car_play/android_auto 0
bluetooth              0
usb_port               0
heated_seats           0
interior_color         0
exterior_color         0
price                  0
dtype: int64

```

```

[27]: print(df1.select_dtypes(include='object').nunique())
categorical_vars = df1.select_dtypes(include='object')
categorical_vars

```

```

brand      25
model     3114
engine     976
transmission 149
fuel_type   7
drivetrain  4
interior_color 904
exterior_color 1841
dtype: int64

```

```

[27]:
      brand      model \
0      Honda  Pilot Special Edition
1  Mercedes-Benz  GL-Class GL 450 4MATIC
2    Mitsubishi  Outlander Sport SE
3      Porsche                968
4      MINI      Hardtop Cooper
...      ...      ...
17463    Land Rover Range Rover Evoque SE
17464    Mitsubishi      Outlander SE
17465      Nissan      GT-R Premium
17466    Maserati      Levante
17467      Kia      Optima Hybrid LX

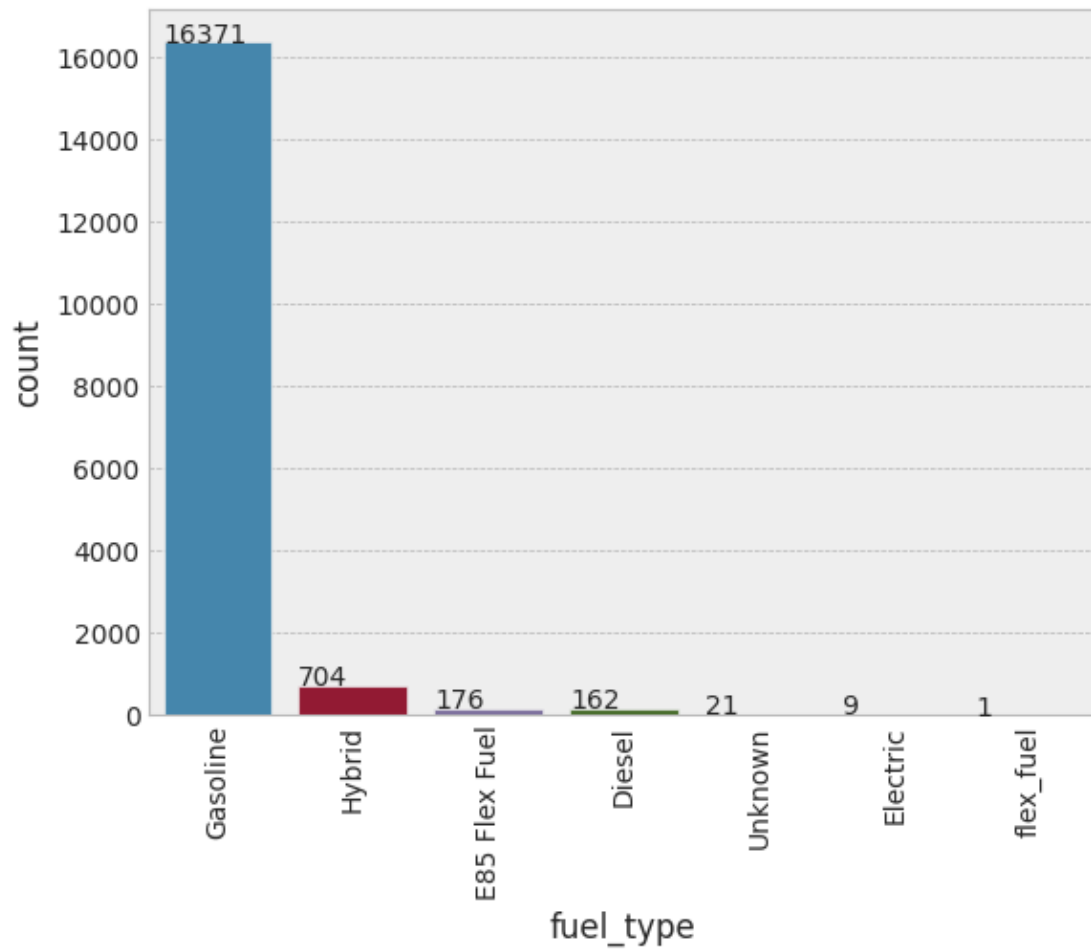
```

	engine	transmission \
0	3.5L V6 24V GDI SOHC	9-Speed Automatic
1	4.7L V8 32V GDI DOHC Twin Turbo	7-Speed Automatic
2	2.0L I4 16V MPFI DOHC	Automatic CVT
3	3.0L I4 16V MPFI DOHC	6-Speed Automatic
4	1.5L I3 12V GDI DOHC Turbo	6-Speed Automatic
...
17463	2.0L I4 16V GDI DOHC Turbo	9-Speed Automatic
17464	2.4L I-4 variable valve control, engine with 1...	Automatic
17465	3.8L V-6 DOHC, variable valve control, twin tu...	Automatic
17466	3.0L V6 24V GDI DOHC Twin Turbo	8-Speed Automatic
17467	2.4L I4 16V MPFI DOHC Hybrid	6-Speed Automatic

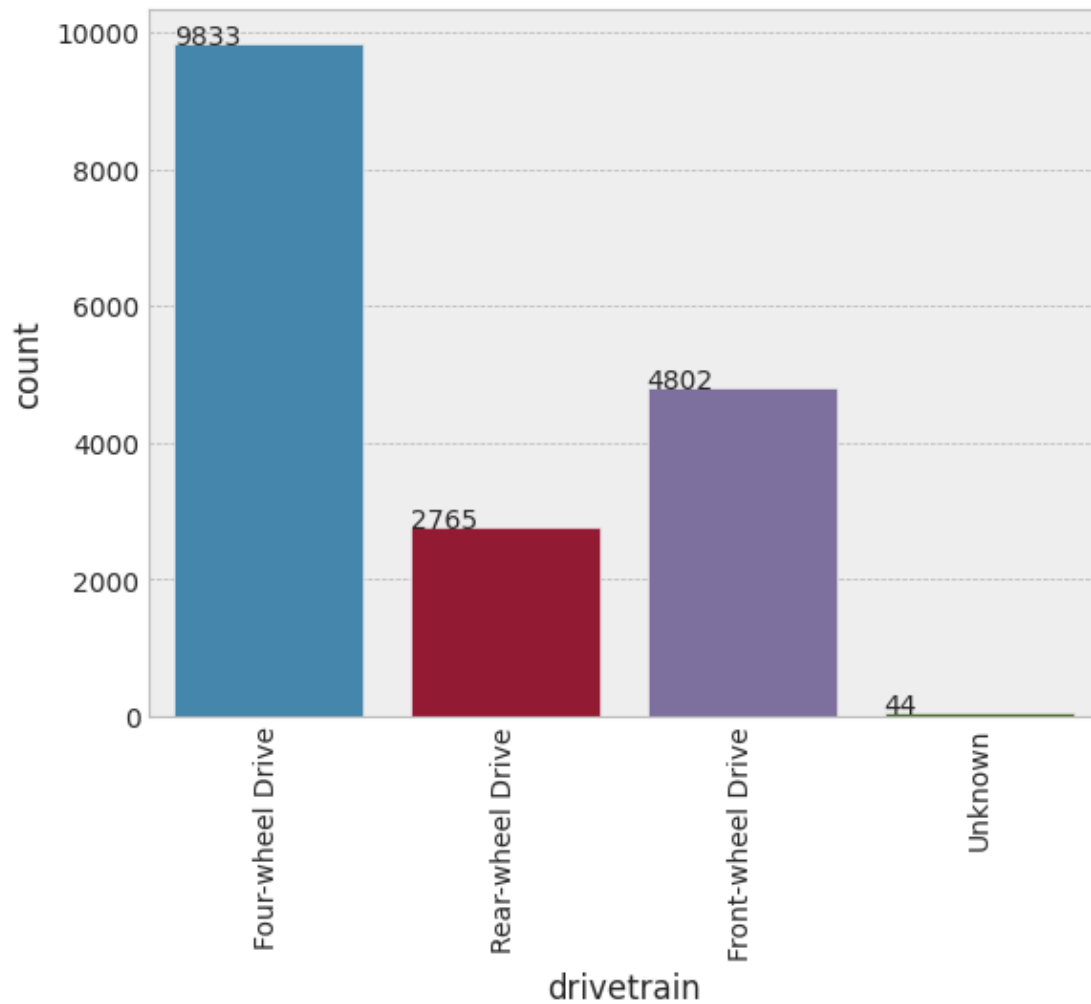
	fuel_type	drivetrain	interior_color	exterior_color
0	Gasoline	Four-wheel Drive	Black	Modern Steel Metallic
1	Gasoline	Four-wheel Drive	Black	Obsidian Black Metallic
2	Gasoline	Four-wheel Drive	Gray	Silver
3	Gasoline	Rear-wheel Drive	Saddle	White
4	Gasoline	Front-wheel Drive	Carbon Black	Electric Blue Metallic
...
17463	Gasoline	Four-wheel Drive	Ebony	Santorini Black
17464	Gasoline	Four-wheel Drive	Black	Mercury Gray Metallic
17465	Gasoline	Four-wheel Drive	Black	Pearl White
17466	Gasoline	Four-wheel Drive	Black	White
17467	Hybrid	Front-wheel Drive	Sand / Black	Black

[17444 rows x 8 columns]

```
[28]: ax = sns.countplot(x='fuel_type', data = categorical_vars)
plt.xticks(rotation=90)
for p in ax.patches:
    ax.annotate('{:.0f}'.format(p.get_height()), (p.get_x(), p.get_height()))
plt.xticks(rotation=90)
sns.set_style('whitegrid')
plt.style.use('bmh')
```



```
[29]: ax = sns.countplot(x='drivetrain', data = categorical_vars)
plt.xticks(rotation=90)
for p in ax.patches:
    ax.annotate('{:.0f}'.format(p.get_height()), (p.get_x(), p.get_height()))
plt.xticks(rotation=90)
sns.set_style('whitegrid')
plt.style.use('bmh')
```



```
[30]: print(df1.select_dtypes(include='float').nunique())
numerical_vars = df1.select_dtypes(include='float')
numerical_vars
```

```
year                44
mileage             16094
engine_size         57
automatic_transmission  2
min_mpg             53
max_mpg             55
damaged             2
first_owner         2
personal_using      2
turbo               2
alloy_wheels        2
adaptive_cruise_control  2
```



```

navigation_system      2
power_liftgate         2
backup_camera          2
keyless_start          2
remote_start           2
sunroof/moonroof       2
automatic_emergency_braking 2
stability_control      2
leather_seats          2
memory_seat            2
third_row_seating      2
apple_car_play/android_auto 2
bluetooth              2
usb_port               2
heated_seats           2
price                  8409
dtype: int64

```

```

[30]:      year  mileage  engine_size  automatic_transmission  min_mpg  \
0      2021.0  51299.0          3.5              1.0         19.0
1      2014.0  79279.0          4.7              1.0         21.0
2      2021.0  31976.0          2.0              1.0         24.0
3      1994.0  92025.0          3.0              1.0         10.0
4      2016.0  43418.0          1.5              1.0         21.0
...      ...      ...      ...      ...      ...
17463  2016.0  106438.0          2.0              1.0         20.0
17464  2019.0  56555.0          2.4              1.0         22.0
17465  2015.0  12091.0          3.8              1.0         17.0
17466  2021.0  31187.0          3.0              1.0         14.0
17467  2013.0  86261.0          2.4              1.0         28.0

      max_mpg  damaged  first_owner  personal_using  turbo  ...  \
0      26.0      0.0          1.0          1.0      0.0  ...
1      28.0      0.0          0.0          1.0      1.0  ...
2      31.0      0.0          1.0          0.0      0.0  ...
3      16.0      1.0          0.0          0.0      0.0  ...
4      31.0      0.0          0.0          0.0      1.0  ...
...      ...      ...      ...      ...      ...
17463  25.0      0.0          0.0          0.0      1.0  ...
17464  30.0      0.0          1.0          1.0      0.0  ...
17465  24.0      0.0          1.0          0.0      0.0  ...
17466  19.0      1.0          0.0          1.0      1.0  ...
17467  37.0      0.0          0.0          1.0      0.0  ...

      automatic_emergency_braking  stability_control  leather_seats  \
0              1.0              1.0              1.0
1              0.0              1.0              1.0

```

2	0.0	1.0	0.0
3	0.0	0.0	1.0
4	0.0	1.0	1.0
...
17463	0.0	1.0	1.0
17464	1.0	1.0	0.0
17465	1.0	1.0	1.0
17466	0.0	1.0	1.0
17467	0.0	1.0	0.0

	memory_seat	third_row_seating	apple_car_play/android_auto	bluetooth	\
0	1.0	1.0	1.0	1.0	
1	1.0	1.0	0.0	1.0	
2	0.0	0.0	0.0	1.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	1.0	
...	
17463	1.0	0.0	0.0	1.0	
17464	0.0	1.0	1.0	1.0	
17465	0.0	0.0	0.0	1.0	
17466	1.0	0.0	0.0	1.0	
17467	0.0	0.0	0.0	1.0	

	usb_port	heated_seats	price
0	0.0	1.0	35999.0
1	0.0	0.0	22349.0
2	0.0	0.0	19499.0
3	0.0	0.0	12960.0
4	1.0	0.0	16500.0
...
17463	1.0	1.0	16495.0
17464	1.0	1.0	23988.0
17465	0.0	1.0	92995.0
17466	0.0	1.0	49850.0
17467	0.0	0.0	15595.0

[17444 rows x 28 columns]

```
[31]: numerical_vars.columns
```

```
[31]: Index(['year', 'mileage', 'engine_size', 'automatic_transmission', 'min_mpg',
        'max_mpg', 'damaged', 'first_owner', 'personal_using', 'turbo',
        'alloy_wheels', 'adaptive_cruise_control', 'navigation_system',
        'power_liftgate', 'backup_camera', 'keyless_start', 'remote_start',
        'sunroof/moonroof', 'automatic_emergency_braking', 'stability_control',
        'leather_seats', 'memory_seat', 'third_row_seating',
        'apple_car_play/android_auto', 'bluetooth', 'usb_port', 'heated_seats',
```

```

    'price'],
    dtype='object')

```

```

[32]: num_var = numerical_vars[['damaged', 'first_owner', 'personal_using', 'turbo',
    'alloy_wheels', 'adaptive_cruise_control', 'navigation_system',
    'power_liftgate', 'backup_camera', 'keyless_start', 'remote_start',
    'sunroof/moonroof', 'automatic_emergency_braking', 'stability_control',
    'leather_seats', 'memory_seat', 'third_row_seating',
    'apple_car_play/android_auto', 'bluetooth', 'usb_port', 'heated_seats']]
num_var

```

```

[32]:
      damaged  first_owner  personal_using  turbo  alloy_wheels  \
0          0.0          1.0          1.0    0.0          1.0
1          0.0          0.0          1.0    1.0          1.0
2          0.0          1.0          0.0    0.0          0.0
3          1.0          0.0          0.0    0.0          1.0
4          0.0          0.0          0.0    1.0          1.0
...
17463      0.0          0.0          0.0    1.0          1.0
17464      0.0          1.0          1.0    0.0          1.0
17465      0.0          1.0          0.0    0.0          1.0
17466      1.0          0.0          1.0    1.0          1.0
17467      0.0          0.0          1.0    0.0          1.0

```

```

      adaptive_cruise_control  navigation_system  power_liftgate  \
0                          1.0              0.0              1.0
1                          0.0              1.0              1.0
2                          0.0              0.0              0.0
3                          0.0              0.0              0.0
4                          0.0              0.0              0.0
...
17463                      0.0              0.0              0.0
17464                      0.0              0.0              0.0
17465                      0.0              1.0              0.0
17466                      0.0              1.0              1.0
17467                      0.0              0.0              0.0

```

```

      backup_camera  keyless_start  ...  sunroof/moonroof  \
0              1.0          0.0  ...              1.0
1              1.0          0.0  ...              1.0
2              1.0          0.0  ...              0.0
3              0.0          0.0  ...              0.0
4              0.0          1.0  ...              0.0
...
17463          1.0          1.0  ...              0.0
17464          1.0          1.0  ...              0.0
17465          1.0          1.0  ...              0.0

```

17466	1.0	1.0	...	1.0
17467	0.0	1.0	...	0.0

	automatic_emergency_braking	stability_control	leather_seats	\
0	1.0	1.0	1.0	
1	0.0	1.0	1.0	
2	0.0	1.0	0.0	
3	0.0	0.0	1.0	
4	0.0	1.0	1.0	
...	
17463	0.0	1.0	1.0	
17464	1.0	1.0	0.0	
17465	1.0	1.0	1.0	
17466	0.0	1.0	1.0	
17467	0.0	1.0	0.0	

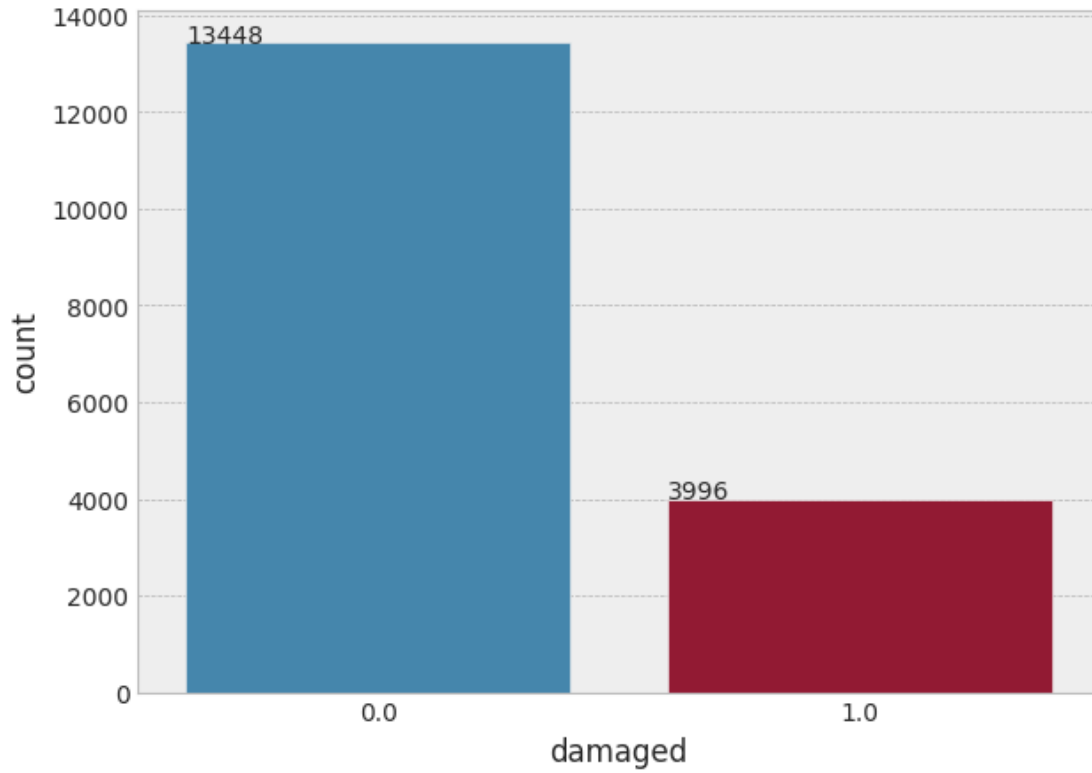
	memory_seat	third_row_seating	apple_car_play/android_auto	bluetooth	\
0	1.0	1.0	1.0	1.0	
1	1.0	1.0	0.0	1.0	
2	0.0	0.0	0.0	1.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	1.0	
...	
17463	1.0	0.0	0.0	1.0	
17464	0.0	1.0	1.0	1.0	
17465	0.0	0.0	0.0	1.0	
17466	1.0	0.0	0.0	1.0	
17467	0.0	0.0	0.0	1.0	

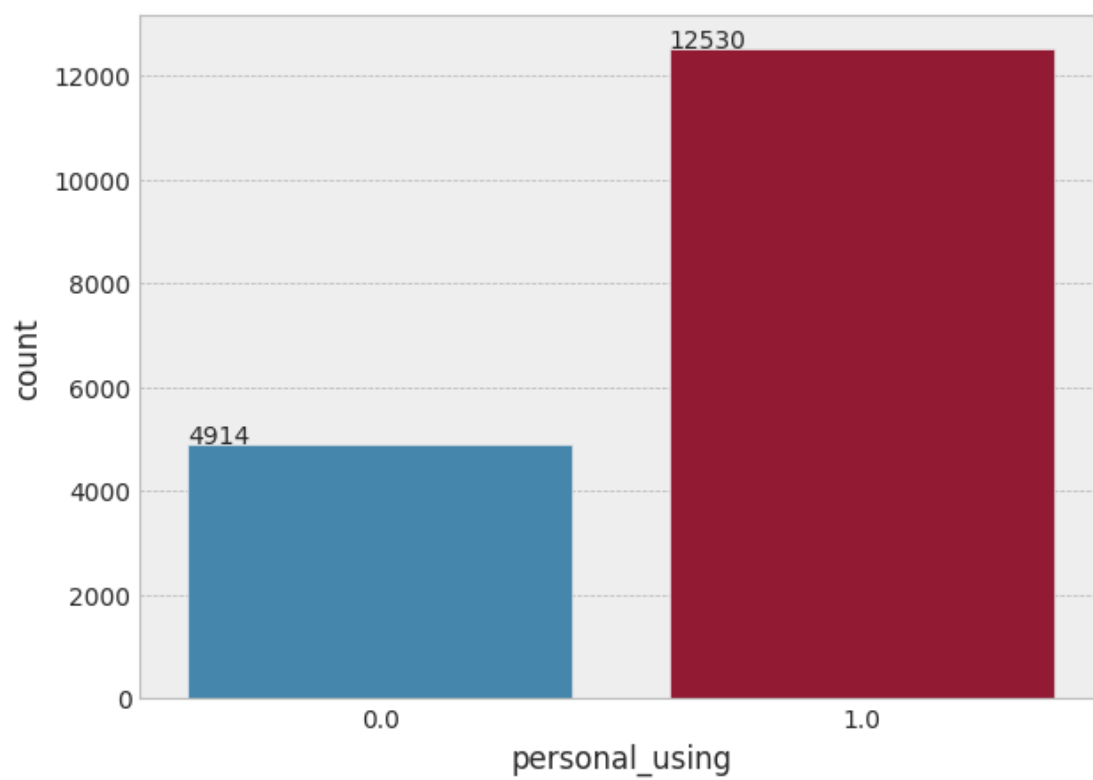
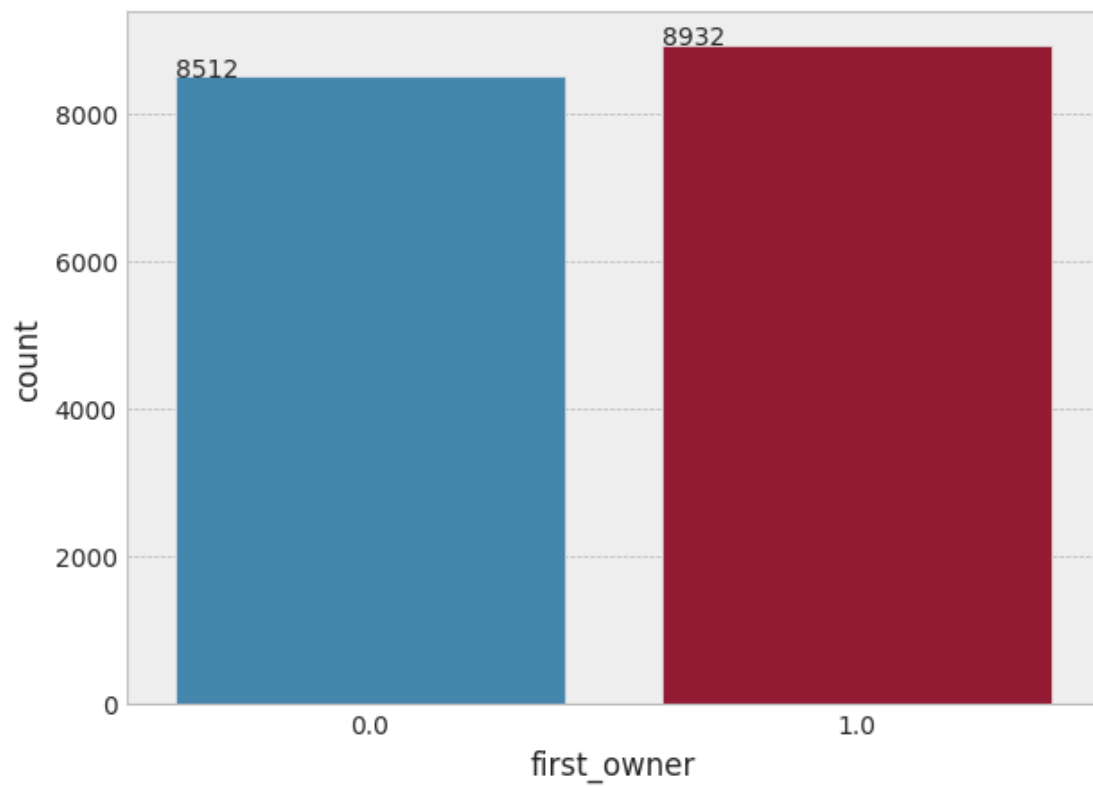
	usb_port	heated_seats
0	0.0	1.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	1.0	0.0
...
17463	1.0	1.0
17464	1.0	1.0
17465	0.0	1.0
17466	0.0	1.0
17467	0.0	0.0

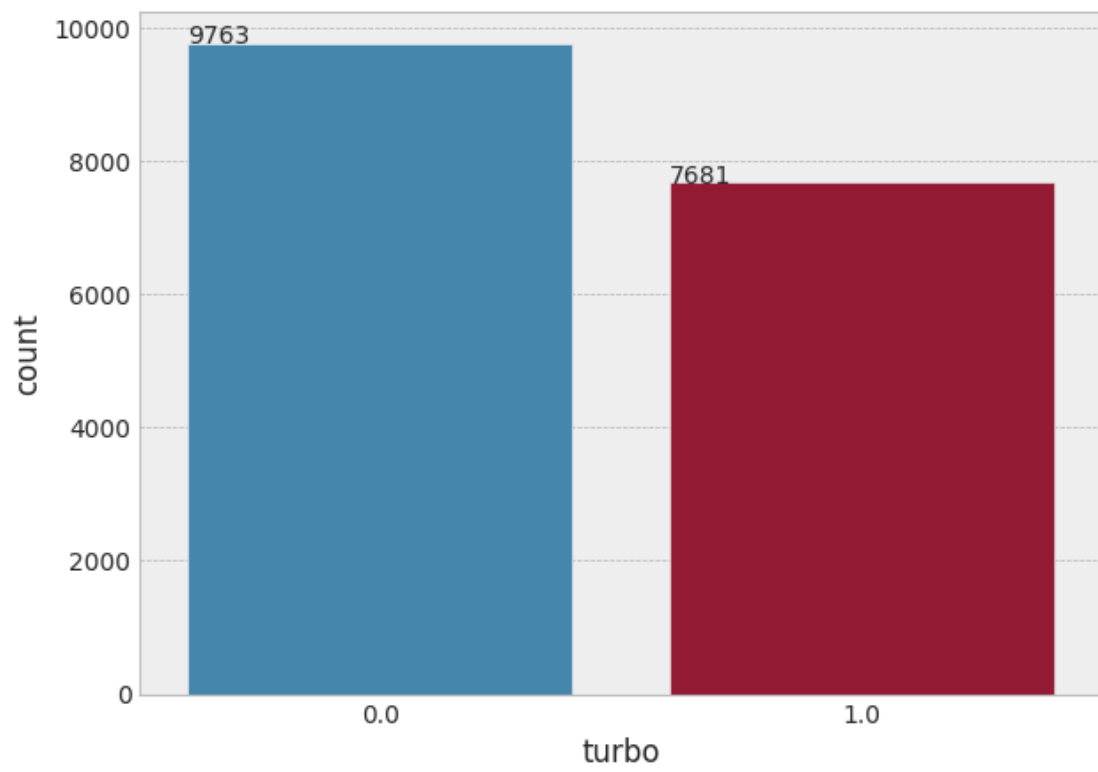
[17444 rows x 21 columns]

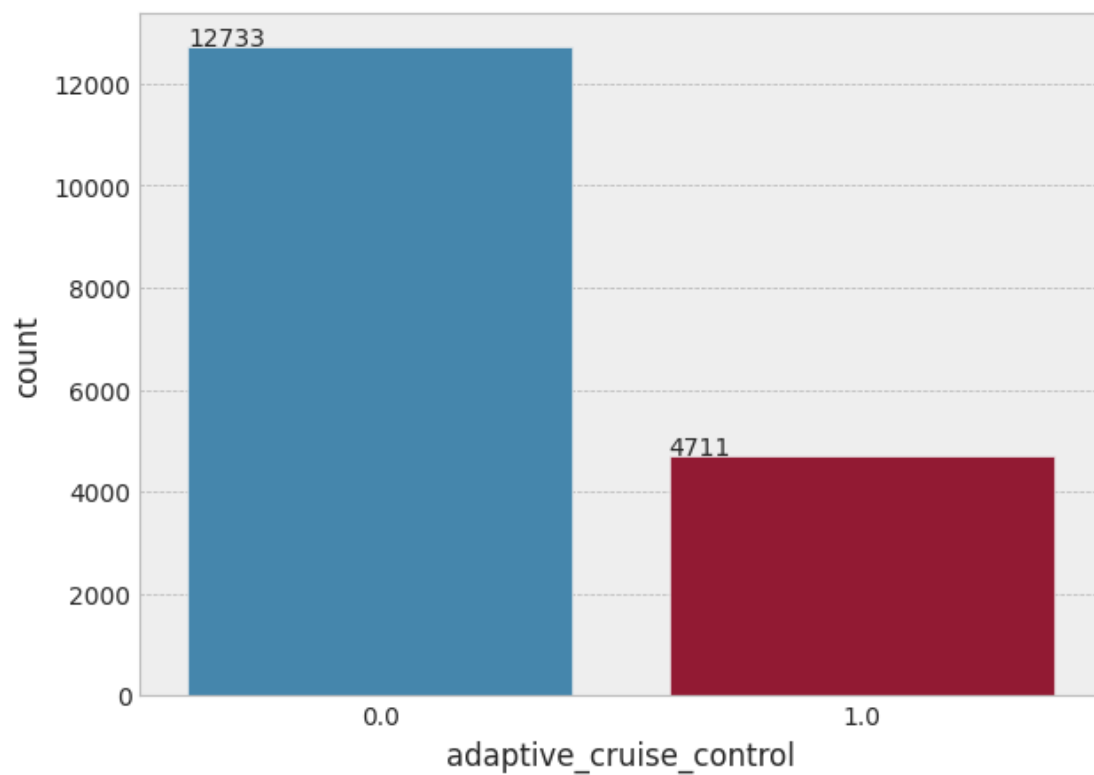
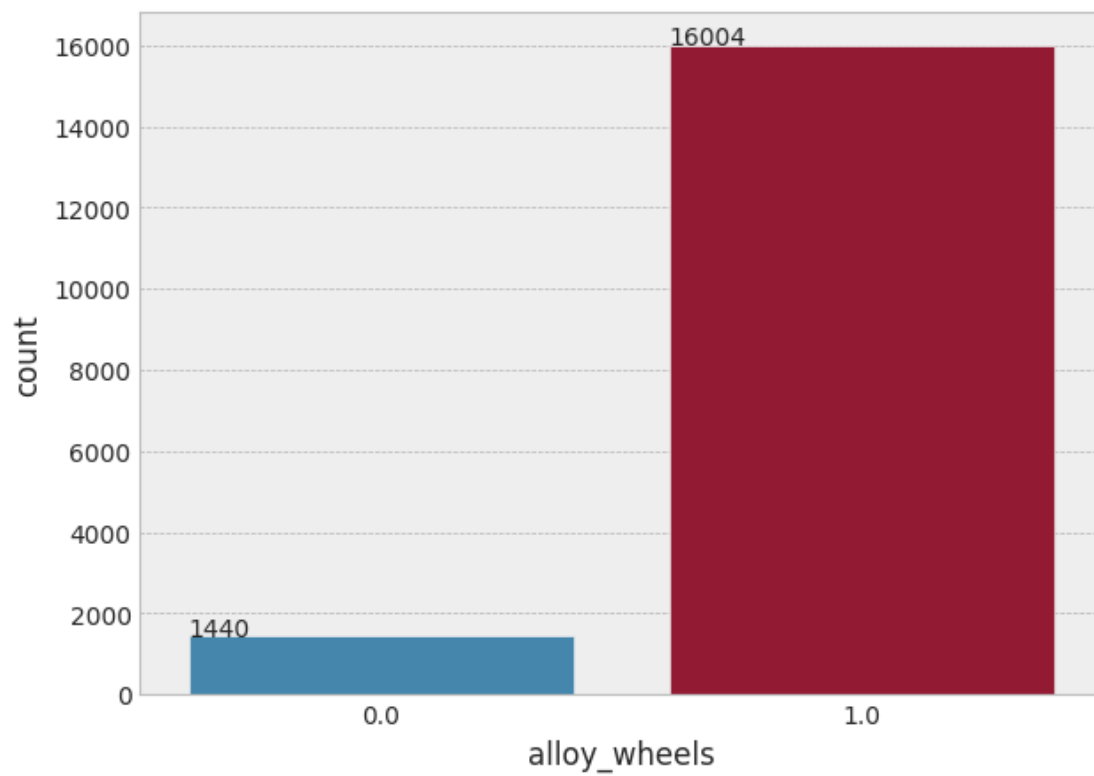
```
[33]: for i in num_var:
      plt.figure(figsize=(7,5))
      ax = sns.countplot(x=num_var[i], data=num_var)
```

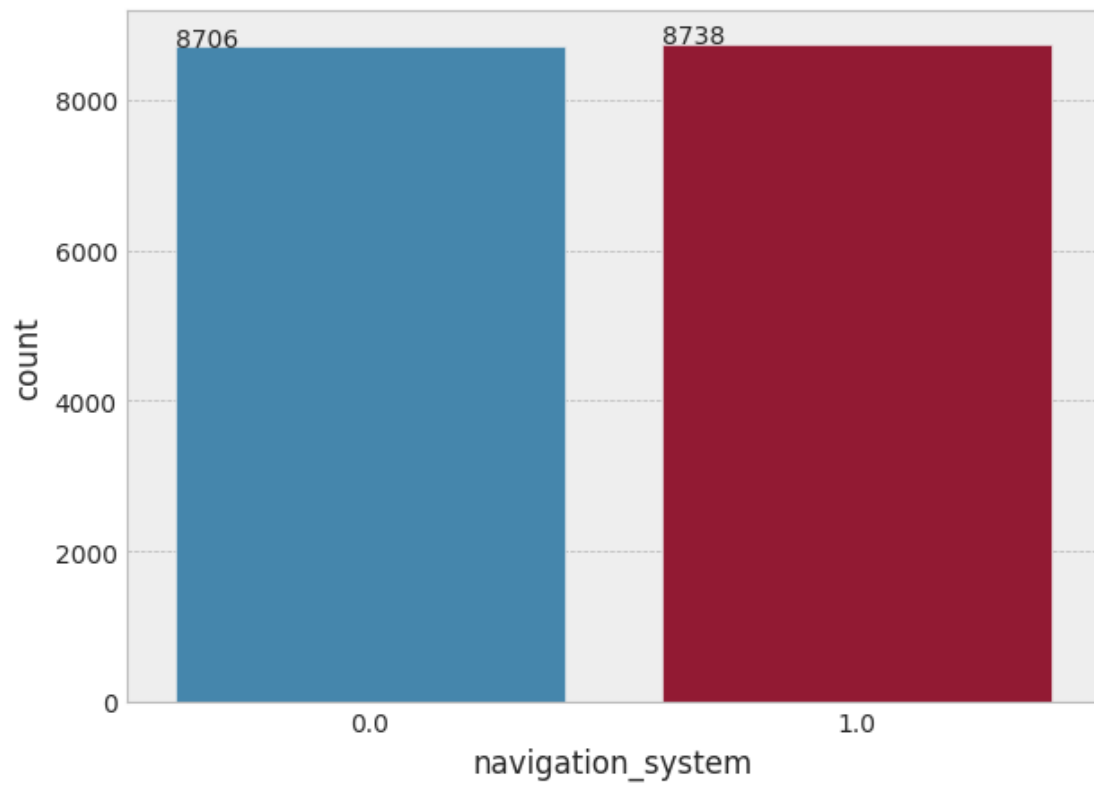
```
for p in ax.patches:  
    ax.annotate('{:.0f}'.format(p.get_height()), (p.get_x(), p.get_height()))  
sns.set_style('whitegrid')  
plt.style.use('bmh')
```

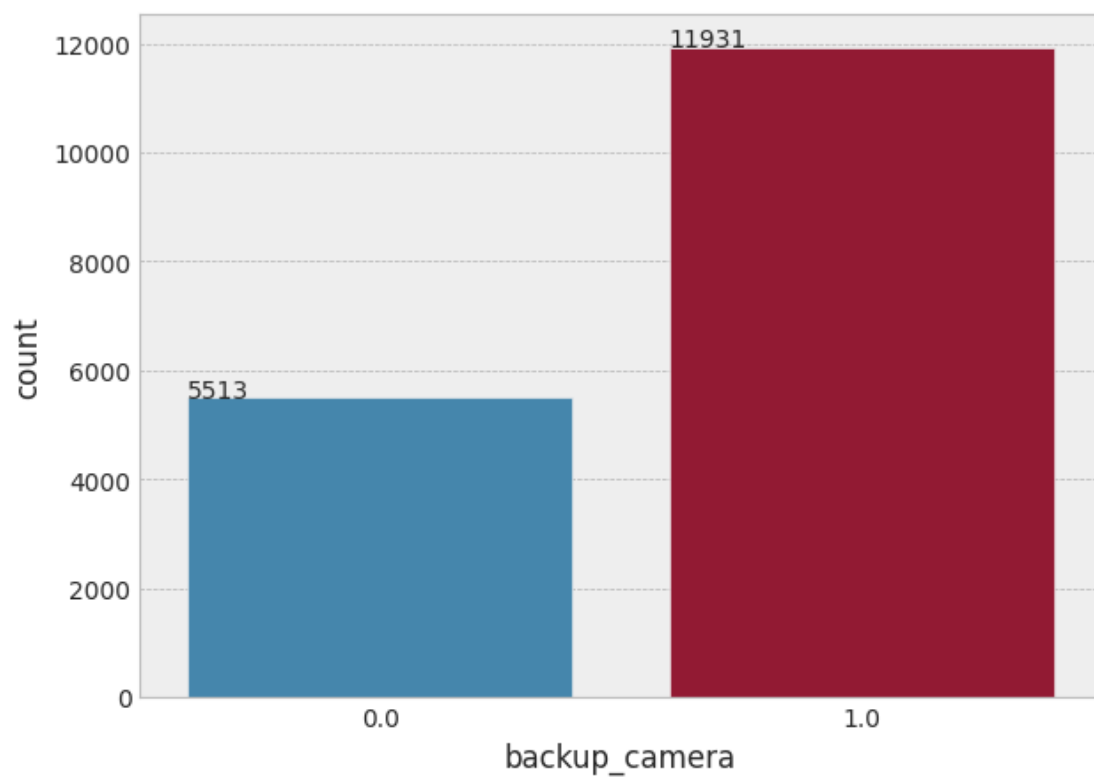
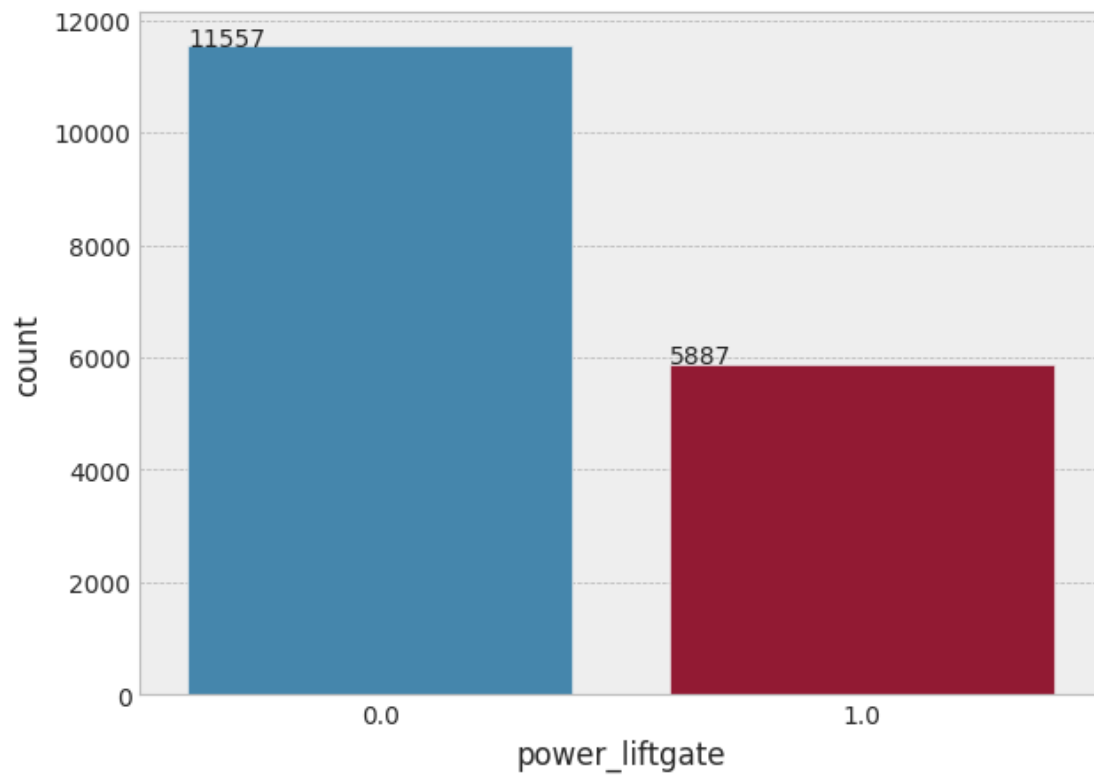


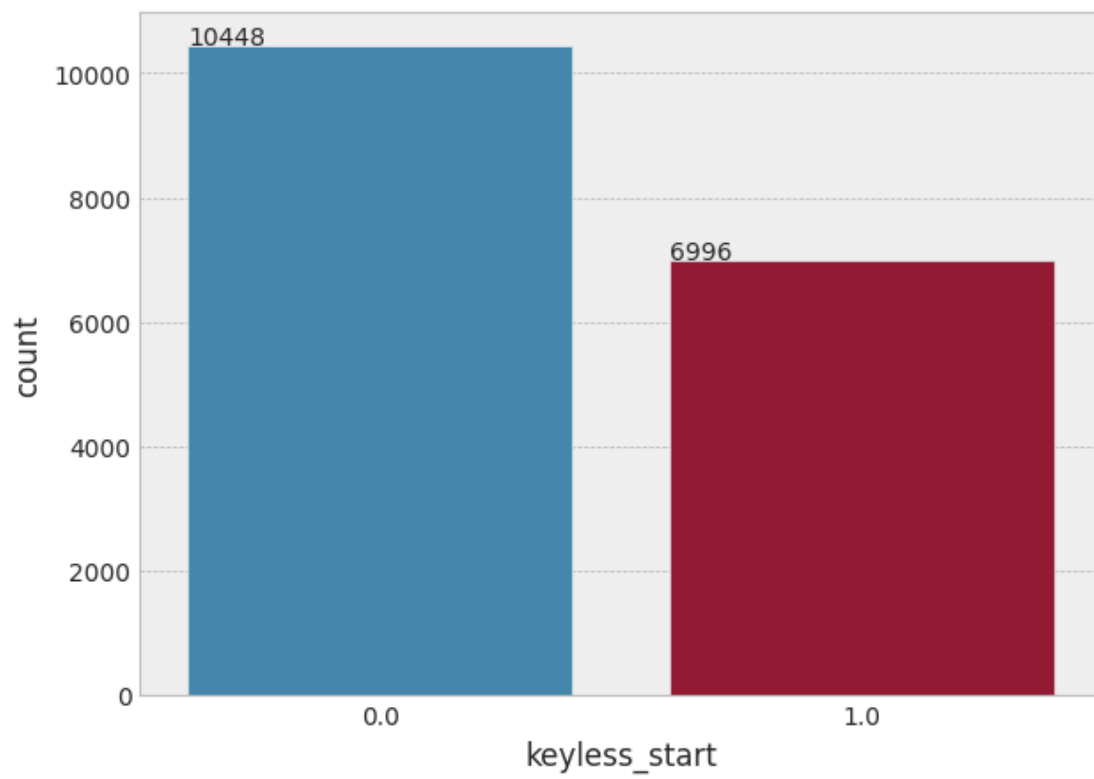


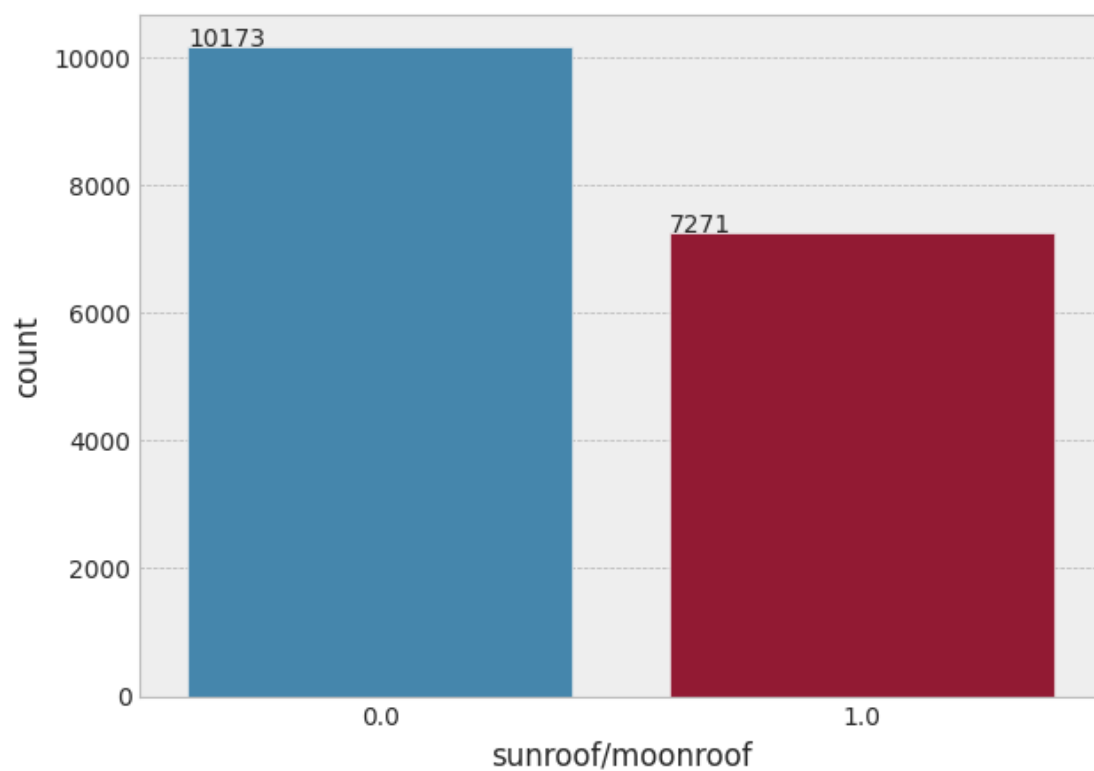
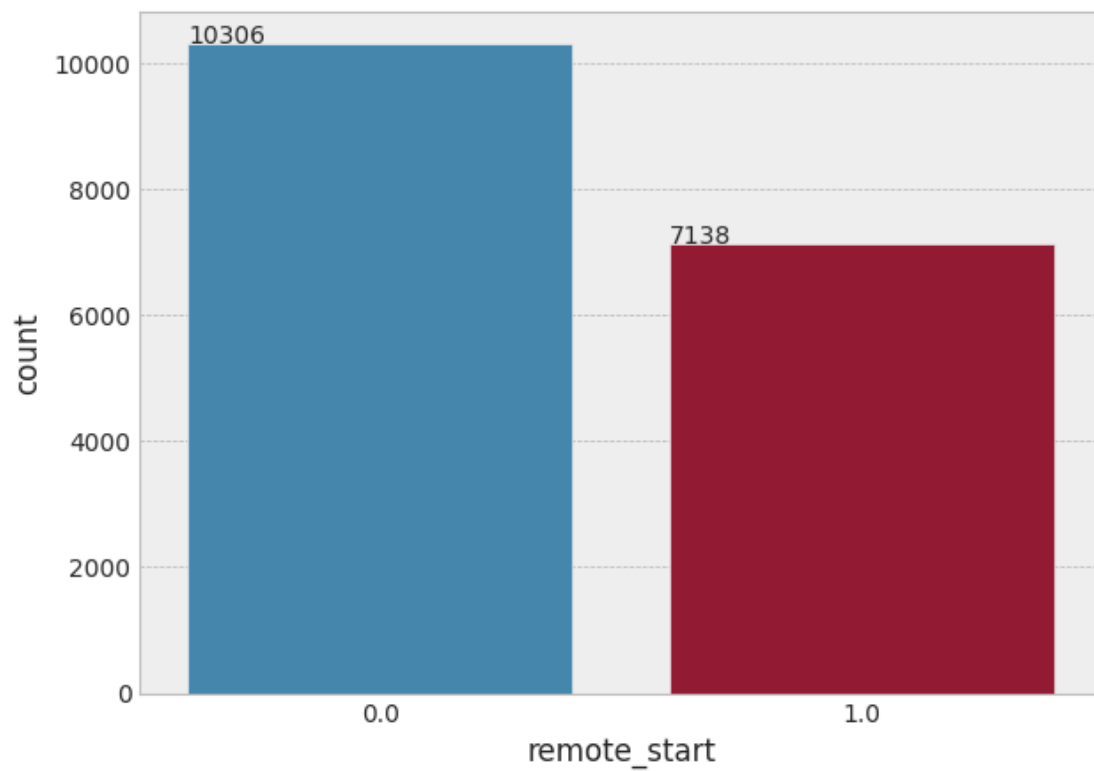


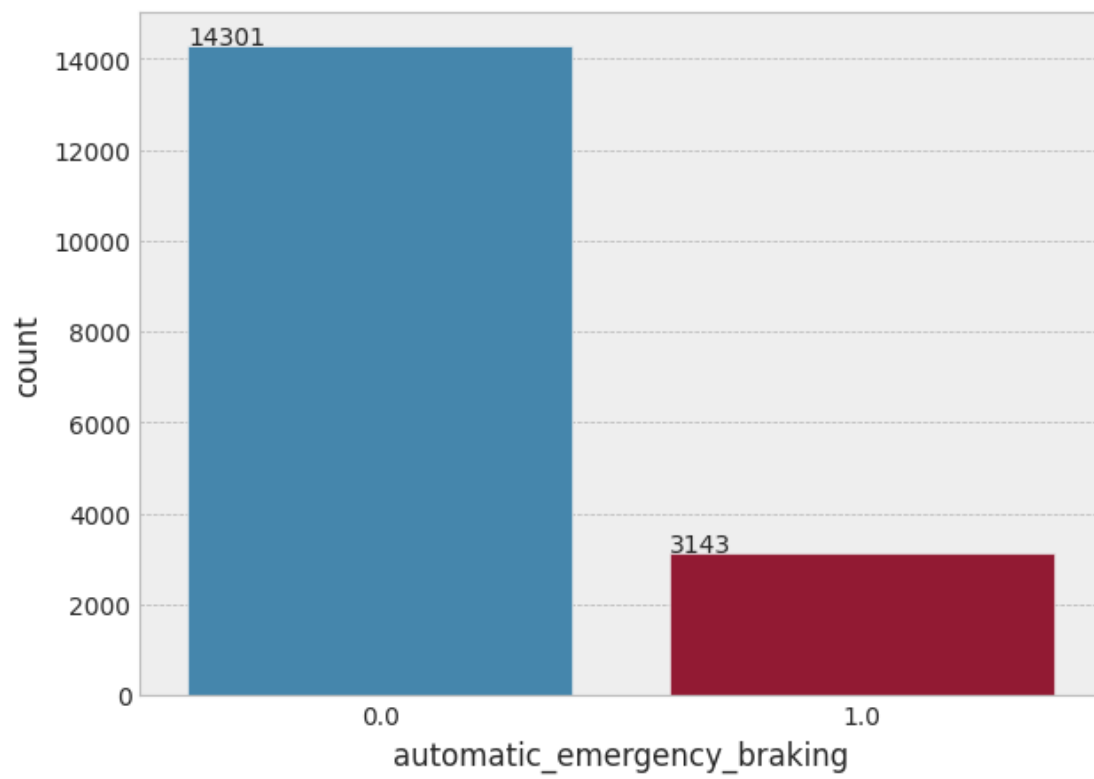


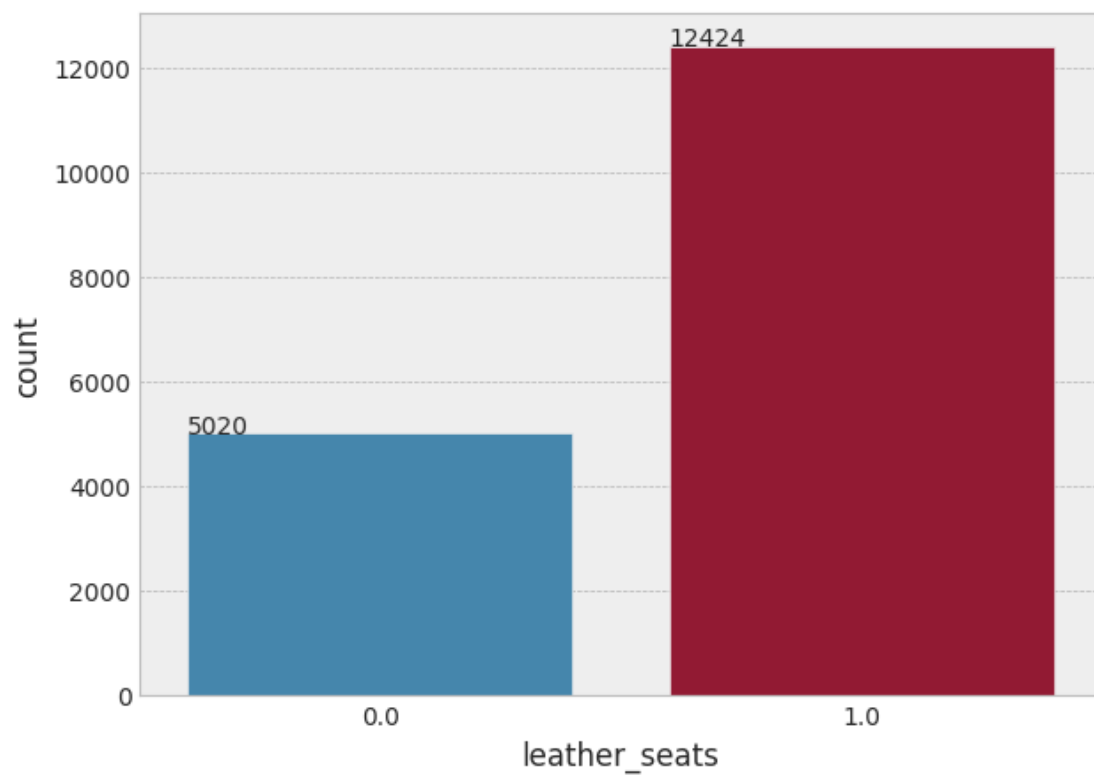
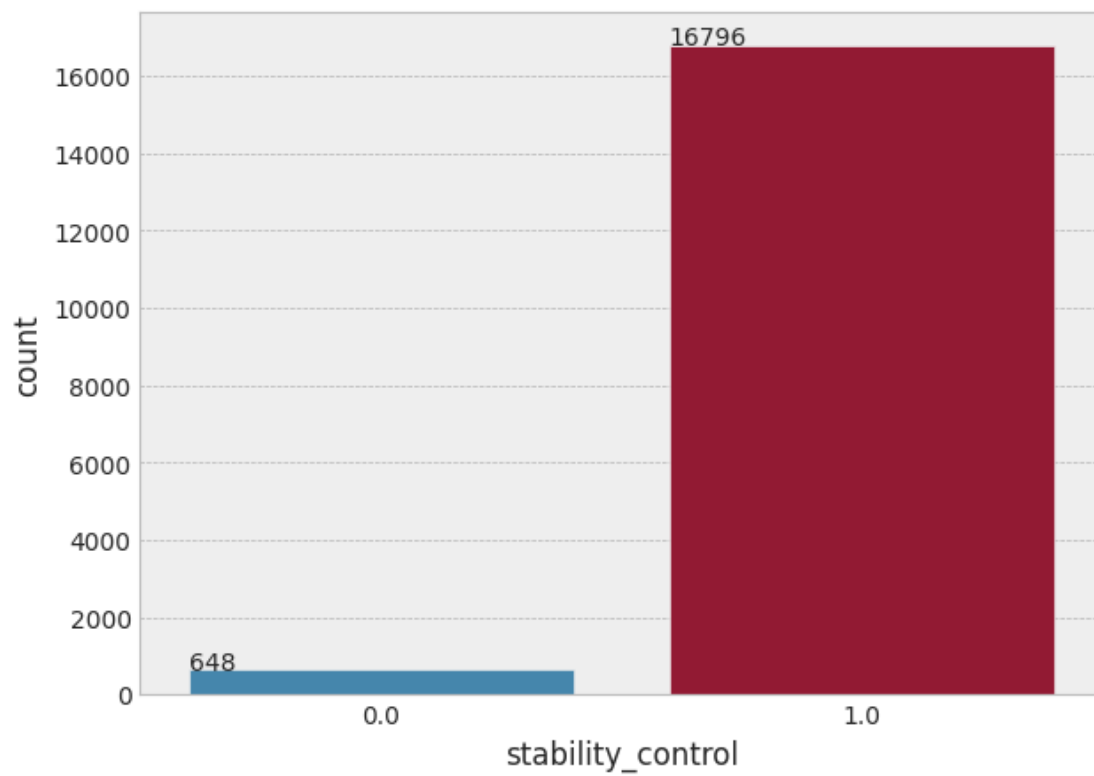


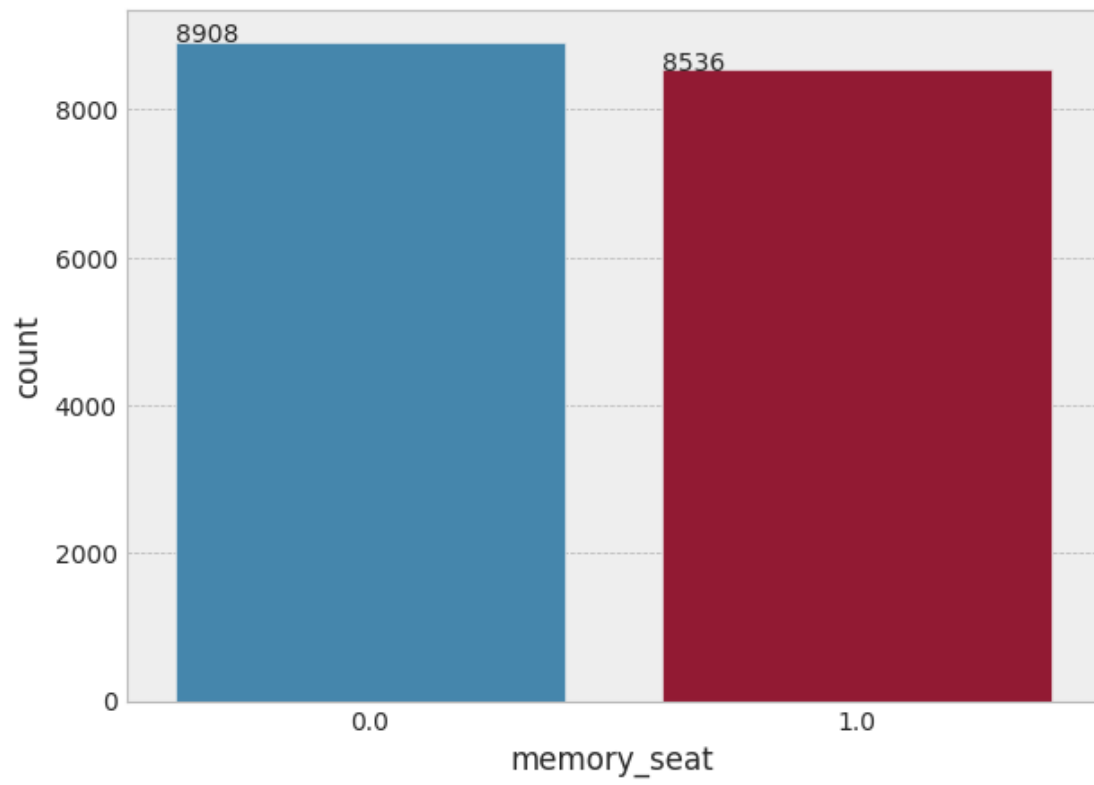


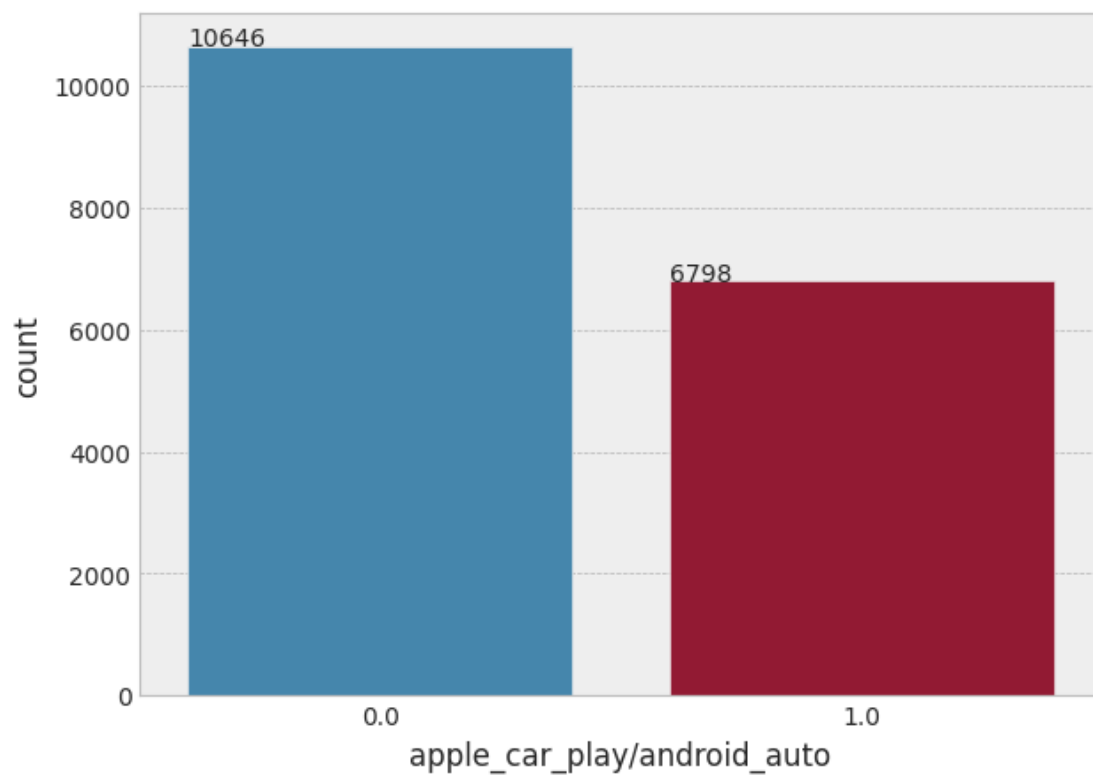
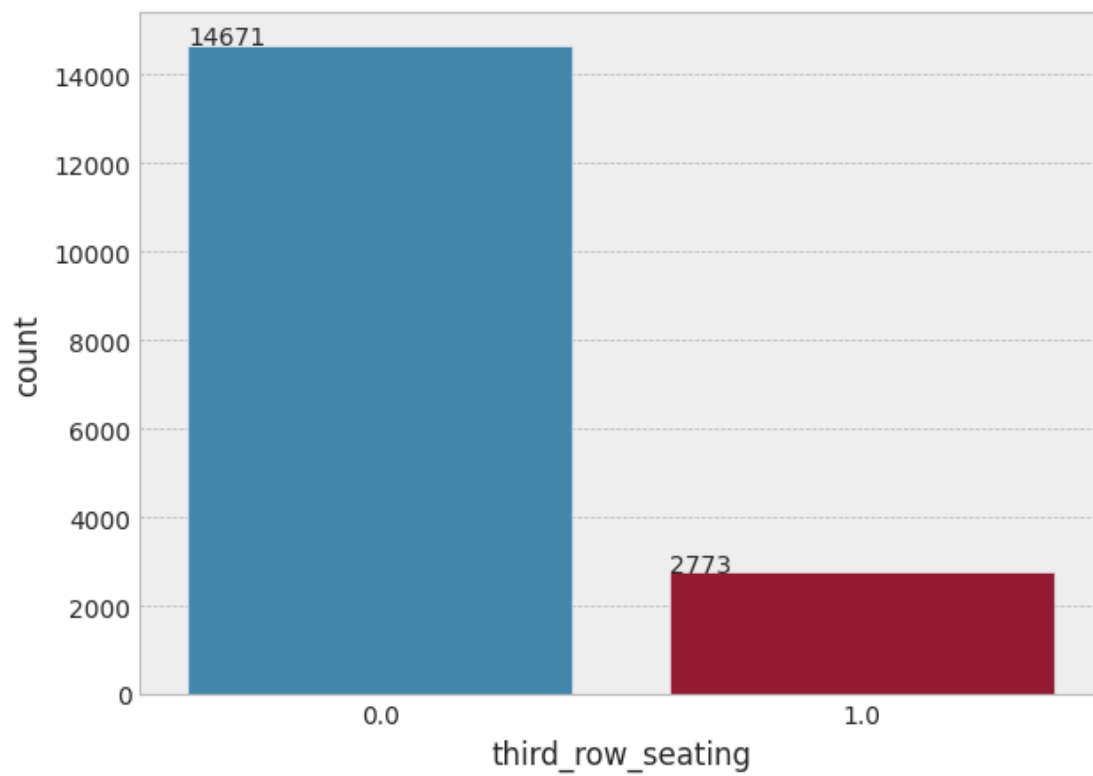


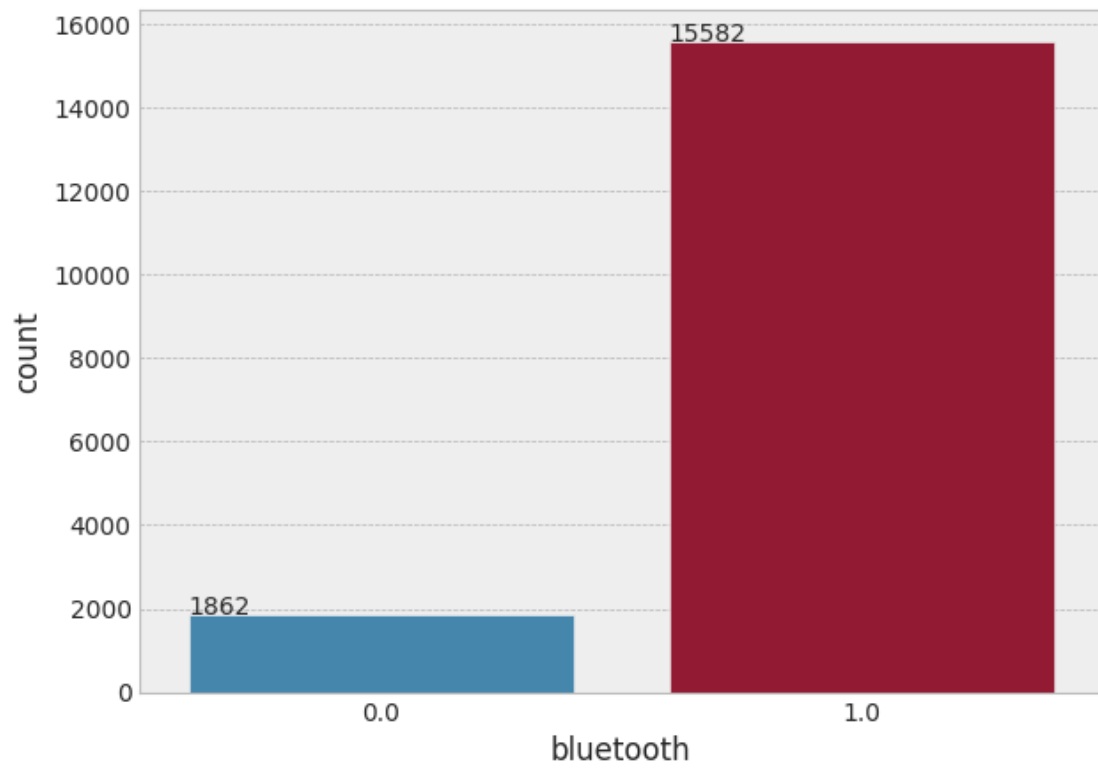


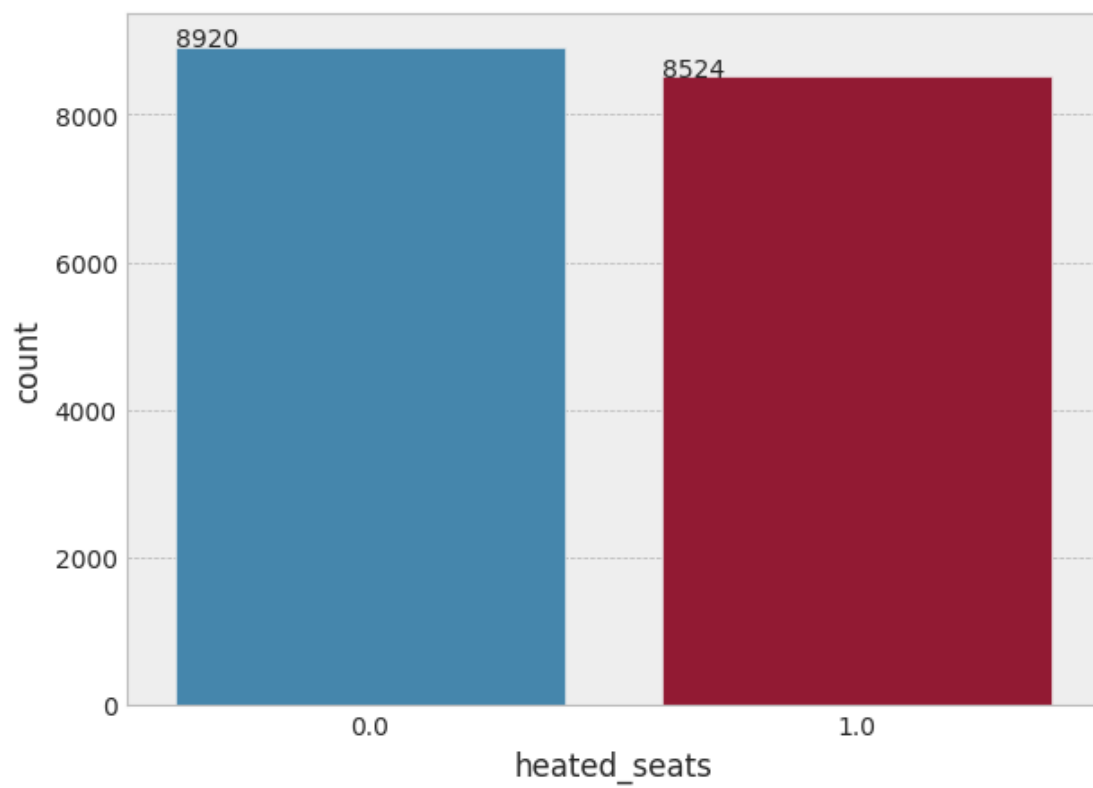
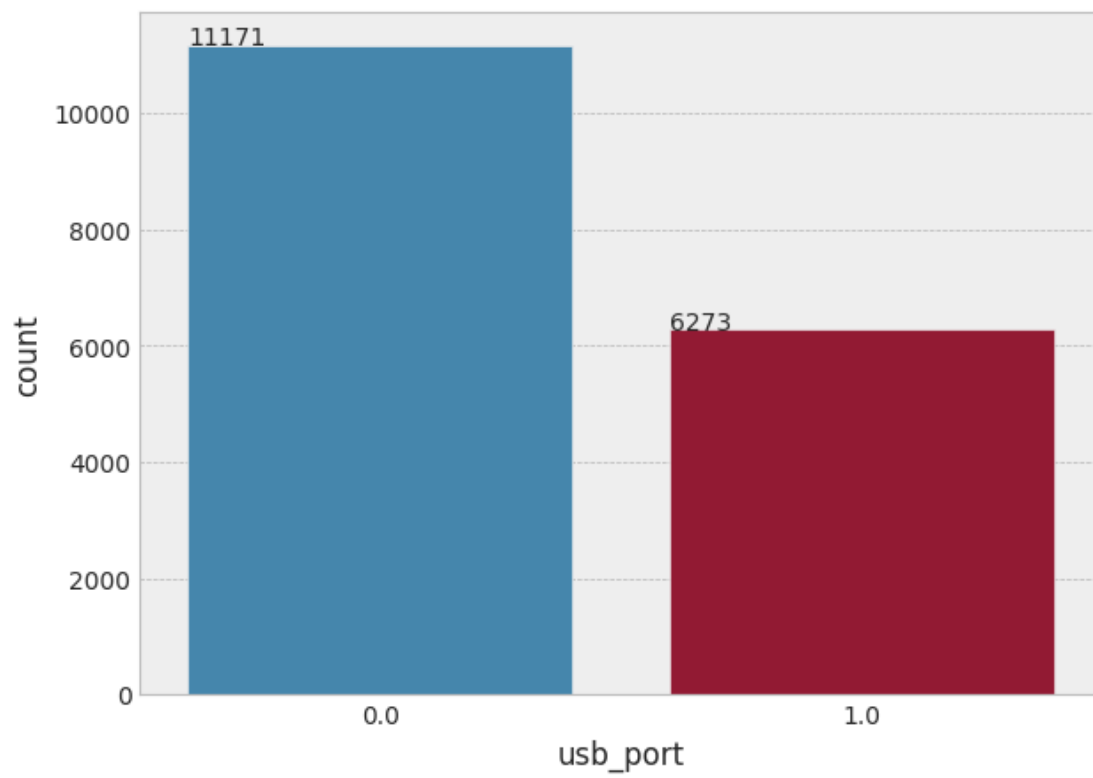












8 Label Encoding for the Object Datatypes

```
[34]: #Splitting into dependent and independent variable
x = df1.drop('price', axis=1)
y = df1['price']
```

8.1 Label Encoding

```
[35]: import sklearn
from sklearn import preprocessing
from sklearn.preprocessing import LabelEncoder

#looping over each columns in the dataframe where datatypes is object
for i in x.select_dtypes(include=['object']).columns:

    #Initializing the labelencoder
    encoder = LabelEncoder()

    #Fitting the label encoder
    encoder.fit(x[i].unique())

    #Transformation
    x[i] = encoder.transform(x[i])

#print the df
x.head()
```

```
[35]:   brand  model   year  mileage  engine  engine_size  transmission  \
0      7   1905  2021.0  51299.0    452         3.5           84
1     17   1175  2014.0  79279.0    614         4.7           57
2     18   1823  2021.0  31976.0    150         2.0          109
3     20    198  1994.0  92025.0    351         3.0           32
4     14   1374  2016.0  43418.0     40         1.5           32

   automatic_transmission  fuel_type  drivetrain  ...  stability_control  \
0                      1.0         3           0  ...                1.0
1                      1.0         3           0  ...                1.0
2                      1.0         3           0  ...                1.0
3                      1.0         3           2  ...                0.0
4                      1.0         3           1  ...                1.0

   leather_seats  memory_seat  third_row_seating  apple_car_play/android_auto  \
```

0	1.0	1.0	1.0	1.0
1	1.0	1.0	1.0	0.0
2	0.0	0.0	0.0	0.0
3	1.0	0.0	0.0	0.0
4	1.0	0.0	0.0	0.0

	bluetooth	usb_port	heated_seats	interior_color	exterior_color
0	1.0	0.0	1.0	77	1052
1	1.0	0.0	0.0	77	1154
2	1.0	0.0	0.0	428	1496
3	0.0	0.0	0.0	736	1766
4	1.0	1.0	0.0	206	598

[5 rows x 35 columns]

8.2 Feature Scaling

```
[36]: #Feature Scaling
from sklearn.preprocessing import StandardScaler
scale = StandardScaler()
x = scale.fit_transform(x)
x = pd.DataFrame(x, columns=df1.drop('price', axis=1).columns)
x.head()
```

```
[36]:      brand      model      year  mileage  engine  engine_size  \
0 -0.624677  0.447198  0.738608  0.053990  0.546656    0.128544
1  0.771964 -0.380692 -0.874548  0.805256  1.273838    0.386823
2  0.911628  0.354202  0.738608 -0.464834 -0.808957   -0.194303
3  1.190956 -1.488703 -5.483563  1.147488  0.093289    0.020928
4  0.352971 -0.155007 -0.413646 -0.157615 -1.302723   -0.301919

      transmission  automatic_transmission  fuel_type  drivetrain  ...  \
0      0.598064                0.28811    0.013621   -0.793025  ...
1     -0.280975                0.28811    0.013621   -0.793025  ...
2      1.411989                0.28811    0.013621   -0.793025  ...
3     -1.094900                0.28811    0.013621    1.850998  ...
4     -1.094900                0.28811    0.013621    0.528986  ...

      stability_control  leather_seats  memory_seat  third_row_seating  \
0      0.196419      0.635655      1.021558      2.300143
1      0.196419      0.635655      1.021558      2.300143
2      0.196419     -1.573182     -0.978897     -0.434755
3     -5.091145      0.635655     -0.978897     -0.434755
4      0.196419      0.635655     -0.978897     -0.434755

      apple_car_play/android_auto  bluetooth  usb_port  heated_seats  \
0      1.251419      0.345683 -0.749362      1.022965
```

1	-0.799093	0.345683	-0.749362	-0.977551
2	-0.799093	0.345683	-0.749362	-0.977551
3	-0.799093	-2.892822	-0.749362	-0.977551
4	-0.799093	0.345683	1.334469	-0.977551


	interior_color	exterior_color
0	-0.810343	0.279443
1	-0.810343	0.462772
2	0.550389	1.077463
3	1.744422	1.562745
4	-0.310245	-0.536551

[5 rows x 35 columns]

9 Correlation Heatmap

```
[37]: plt.figure(figsize=(20,16))
      corr = df1.corr()
      sns.heatmap(corr, vmin=-1, vmax=1, cmap='Blues',annot=True)
```

```
[37]: <Axes: >
```



10 Train Test Split

```
[38]: #Train set will be 80% while the test set will be 20%
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

metrics_score = []

#Spltting of the set
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
    ↪random_state = 0)

#Checking the shape
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
```

```
print(y_test.shape)
```

```
(13955, 35)
(3489, 35)
(13955,)
(3489,)
```

10.1 Decision Tree Regressor

```
[39]: #import the libraries
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import GridSearchCV

dtree = DecisionTreeRegressor()

#Define the hyperparameters
param_grid = {
    'max_depth': [2,4,6,8],
    'min_samples_split': [2,4,6,8],
    'min_samples_leaf': [1,2,3,4],
    'max_features': ['auto', 'sqrt', 'log2']
}

#Lets create the GridSearchCV object
grid_search = GridSearchCV(dtree, param_grid, cv=5,
    ↪scoring='neg_mean_squared_error')

#Fitting
grid_search.fit(x_train, y_train)

#Checking the best hyperparameters
grid_search.best_params_
```

```
[39]: {'max_depth': 8,
      'max_features': 'auto',
      'min_samples_leaf': 4,
      'min_samples_split': 6}
```

```
[40]: from sklearn.tree import DecisionTreeRegressor
dtree = DecisionTreeRegressor(random_state= 0, max_depth=8,
    ↪max_features='auto', min_samples_leaf=4, min_samples_split=4)
dtree.fit(x_train,y_train)
```

```
[40]: DecisionTreeRegressor(max_depth=8, max_features='auto', min_samples_leaf=4,
    min_samples_split=4, random_state=0)
```

```
[41]: y_pred = dtree.predict(x_test)
      y_pred
```

```
[41]: array([15845.15384615, 27397.48081841, 29689.2739726 , ...,
          35570.91262136, 25444.63513514, 43612.4      ])
```

```
[42]: from sklearn import metrics
      import math
      dt_mse = metrics.mean_squared_error(y_test,y_pred)
      dt_mae = metrics.mean_absolute_error(y_test,y_pred)
      dt_r2 = metrics.r2_score(y_test,y_pred)
      dt_rmse = math.sqrt(dt_mse)

      print('MSE is {}'.format(dt_mse))
      print('MAE is {}'.format(dt_mae))
      print('R2 is {}'.format(dt_r2))
      print('RMSE is {}'.format(dt_rmse))

      metrics_score.append(dt_mse * 100)
      metrics_score.append(dt_mae * 100)
      metrics_score.append(dt_r2)
      metrics_score.append(dt_rmse * 100)
```

```
MSE is 438161061.9192889
MAE is 7068.975484154903
R2 is 0.43445646109272107
RMSE is 20932.297100874737
```

11 Feature Importance

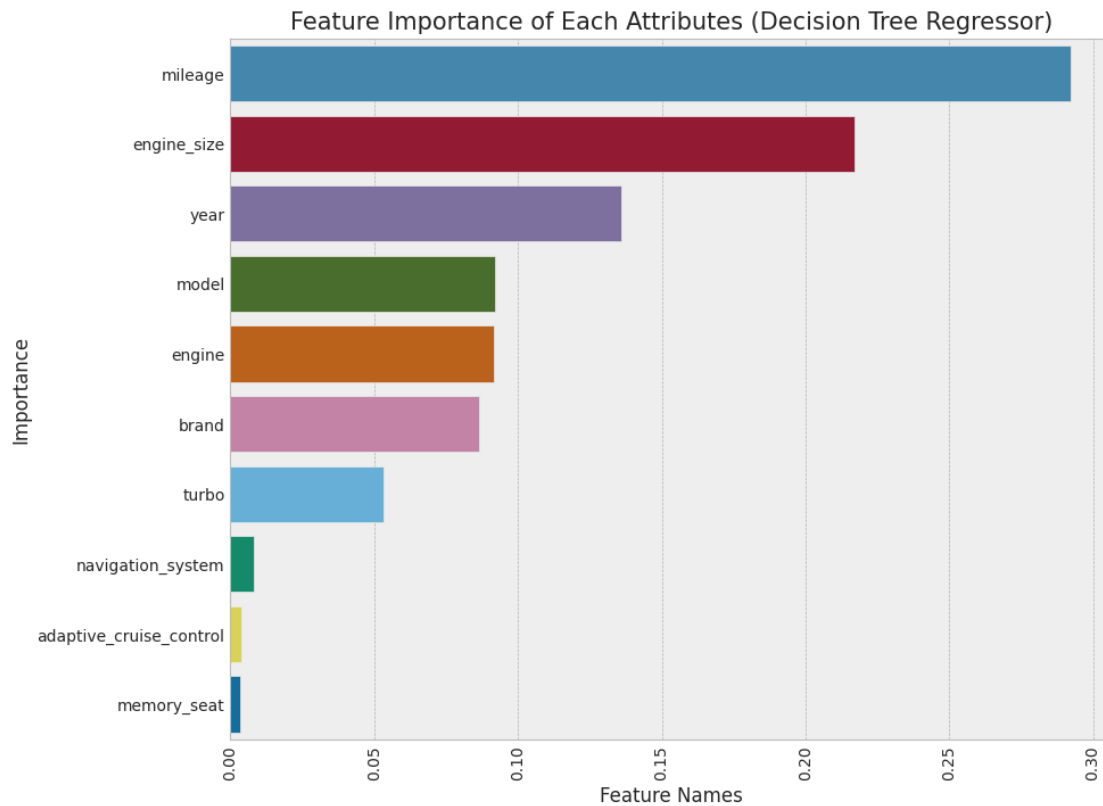
```
[43]: imp_df = pd.DataFrame({
      'Feature Name': pd.DataFrame(x_train).columns,
      'Importance': dtree.feature_importances_
      })

      fi = imp_df.sort_values(by='Importance', ascending=False)

      fi2 = fi.head(10)
      plt.figure(figsize=(10,8))
      sns.barplot(data=fi2, x = 'Importance',y='Feature Name')
      plt.title('Feature Importance of Each Attributes (Decision Tree_
      ↪Regressor)',fontsize=15)
      plt.xlabel('Feature Names')
      plt.xticks(rotation=90)
      plt.ylabel('Importance')
```



```
plt.show()
```



12 Random Forest Regressor

```
[44]: # import the libraries
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV

rf = RandomForestRegressor()
```

```
[45]: #Fitting
rf.fit(x_train, y_train)
```

```
[45]: RandomForestRegressor()
```

```
[46]: #Prediction
rf_pred = rf.predict(x_test)
print(rf_pred)
```

```

from sklearn import metrics
import math
rf_mse = metrics.mean_squared_error(y_test, rf_pred)
rf_mae = metrics.mean_absolute_error(y_test, rf_pred)
rf_r2 = metrics.r2_score(y_test, rf_pred)
rf_rmse = math.sqrt(rf_mse)

print('MSE is {}'.format(rf_mse))
print('MAE is {}'.format(rf_mae))
print('R2 is {}'.format(rf_r2))
print('RMSE is {}'.format(rf_rmse))

metrics_score.append(rf_mse * 100)
metrics_score.append(rf_mae * 100)
metrics_score.append(rf_r2)
metrics_score.append(rf_rmse * 100)

```

[34011.92 25562.28 28265.18 ... 34079.08 25436.45 49091.98]

MSE is 405382917.10245174

MAE is 5085.836279736314

R2 is 0.47676389009457987

RMSE is 20134.123201730235

12.1 Gradient Boosting Regressor

[47]: `from sklearn.ensemble import GradientBoostingRegressor`

```

GBR = GradientBoostingRegressor()

#Training
GBR.fit(x_train, y_train)

#Prediction
GBR_pred = GBR.predict(x_test)
print(GBR_pred)

print('***' * 30)
from sklearn import metrics
import math
gbr_mse = metrics.mean_squared_error(y_test, GBR_pred)
gbr_mae = metrics.mean_absolute_error(y_test, GBR_pred)
gbr_r2 = metrics.r2_score(y_test, GBR_pred)
gbr_rmse = math.sqrt(gbr_mse)

print('MSE is {}'.format(gbr_mse))
print('MAE is {}'.format(gbr_mae))
print('R2 is {}'.format(gbr_r2))

```

```

print('RMSE is {}'.format(gbr_rmse))

metrics_score.append(gbr_mse * 100)
metrics_score.append(gbr_mae * 100)
metrics_score.append(gbr_r2)
metrics_score.append(gbr_rmse * 100)

```

```

[57714.46103407 27269.76952061 27208.21224417 ... 36543.23346787
 39998.27828771 48034.6954281 ]
*****
*****
MSE is 305351895.2778703
MAE is 6472.9868286319
R2 is 0.6058759974903891
RMSE is 17474.321024802946

#Accuracy Check

```

12.1.1 Decision Tree Accuracy

```

[69]: dtree_train_acc = dtree.score(x_train, y_train)
print('Training Accuracy :', dtree.score(x_train, y_train))
dtree_test_acc = dtree.score(x_test, y_test)
print('Testing Accuracy :', dtree.score(x_test, y_test))

```

```

Training Accuracy : 0.6400730474433721
Testing Accuracy : 0.43445646109272107

```

12.1.2 Random Forest Accuracy

```

[70]: rf_train_acc = rf.score(x_train, y_train)
print('Training Accuracy :', rf.score(x_train, y_train))
rf_test_acc = rf.score(x_test, y_test)
print('Testing Accuracy :', rf.score(x_test, y_test))

```

```

Training Accuracy : 0.9377186613860334
Testing Accuracy : 0.47676389009457987

```

12.1.3 Gradient Boosting Accuracy

```

[71]: GBR_train_acc = GBR.score(x_train, y_train)
print('Training Accuracy :', GBR.score(x_train, y_train))
GBR_test_acc = GBR.score(x_test, y_test)
print('Testing Accuracy :', GBR.score(x_test, y_test))

```

```

Training Accuracy : 0.8693346667075357
Testing Accuracy : 0.6058759974903891

```

```

[76]: train_score_list = []

train_score_list.append(dtrees_train_acc)
train_score_list.append(rf_train_acc)
train_score_list.append(GBR_train_acc)

train_score_model_list = ['Decision Tree', 'Random Forest', 'Gradient Boosting']

sns.set_style("whitegrid")
sns.color_palette("Paired")
plt.figure(figsize=(7,5))
ax = sns.barplot(y = train_score_list, x = train_score_model_list)
plt.title("Model Training Score")
plt.xlabel('Regression Model')
plt.ylabel("Score of Regression Training Model")
for i in ax.patches:
    width, height=i.get_width(), i.get_height()
    x, y=i.get_xy()
    ax.annotate(f'{round(height,2)}%', (x+width/2, y+height*1.02), ha = 'center')
plt.xticks(rotation=90)

plt.show()

```



```
[77]: test_score_list = []


test_score_list.append(dtree_test_acc)
test_score_list.append(rf_test_acc)
test_score_list.append(GBR_test_acc)

test_score_model_list = ['Decision Tree', 'Random Forest', 'Gradient Boosting']

sns.set_style("whitegrid")
sns.color_palette("Paired")
plt.figure(figsize=(7,5))
ax = sns.barplot(y = test_score_list, x = test_score_model_list)
plt.title("Model Testing Score")
plt.xlabel('Regression Model')
```

```
plt.ylabel("Score of Regression Testing Model")
for i in ax.patches:
    width, height=i.get_width(), i.get_height()
    x, y=i.get_xy()
    ax.annotate(f'{round(height,2)}%', (x+width/2, y+height*1.02), ha='center')
plt.xticks(rotation=90)

plt.show()
```



```
[ ]: !apt-get install texlive texlive-xetex texlive-latex-extra pandoc
      !pip install pypandoc
```

```
[ ]: !jupyter nbconvert--to PDF 'Used Car Listings: Features and Price Prediction.
      ↪ipynb'
```